

# Applied Deep Learning

## Transfer Learning and Object Detection

Mark Schutera

Karlsruhe Institute of Technology | ZF Friedrichshafen AG | Hochschule Ravensburg-Weingarten

1. Deep Learning Foundations

**3. Transfer Learning and Object Detection**

5. Segmentation Networks

8. Deep Reinforcement Learning

10. Generative Adversarial Networks

12. Recurrent Neural Networks

Project proposal

One page on introduction, methods, dataset

Deadline 3. Lecture

### **Intermediate presentation**

Ten minutes on achievements, problems, next steps

Due 7. Lecture

### **Final presentation**

Code and results

Due 14. Lecture

### **Final documentation**

Paper and code on github or jupyter notebook

Deadline 14. Lecture

## Course features

### Sli.do

Every question matters.

Get the app.

Ask questions (with slide number)  
or vote on other students' questions  
during the lecture.

And give direct feedback.

### #TOBEDETERMINED

Questions will be covered  
immediately or in the next lecture in  
more depth.

### Github

Find slides, tutorials, flashcards and  
references on Github.

[https://github.com/schutera/  
DeepLearningLecture](https://github.com/schutera/DeepLearningLecture) [Schutera](#)

You found typos, additional  
material such as links, algorithms,  
papers, literature or want to  
contribute to the slides and lecture  
notes..

..Feel free to contribute, e-mail me.

## Course features

### Sli.do

Every question matters.

Get the app.

Ask questions (with slide number)  
or vote on other students' questions  
during the lecture.

#TOBEDETERMINED

Questions will be covered directly or  
in the next lecture in more depth.

### Github

Typos, additional material such as  
links, algorithms, paper, literature,  
lecture notes..

..Feel free to contribute.

### **Grade Bonus .3**

Prepare flashcards based on Ian  
Goodfellow's Deep Learning Book

- Commit to flashcard set by  
emailing me, first come first serve
- Must be comprehensive

## **Object Detection with neural networks**

Basic Architecture

Overview state-of-the-art approaches

Datasets and benchmarking

Introduction to Transfer Learning


Transfer Learning with neural networks

# Introduction – Object Detection a problem statement


Classification



Classification  
+ Localization



Object Detection



## Introduction – Object Detection a problem statement


*Object detection involves detection and localization of instances of objects from a particular class in a sensor measurement, such as a camera image.*





## Object Detection Approach

1. Establish Bounding Box Hypotheses
2. Generate Features for each Bounding Box
3. Classify and give a confidence for the objects based on the Features




# Neural Network Object Detection - Basic structure

## Object Detection Approach

1. Establish Bounding Box Hypotheses
2. Generate Features for each Bounding Box
3. Classify and give a confidence for the objects based on the Features

*This can also be done with conventional approaches, such as sliding windows combined with classification.*



## YOLO: You Only Look Once

Discretization of the input image into a  $S \times S$  grid.

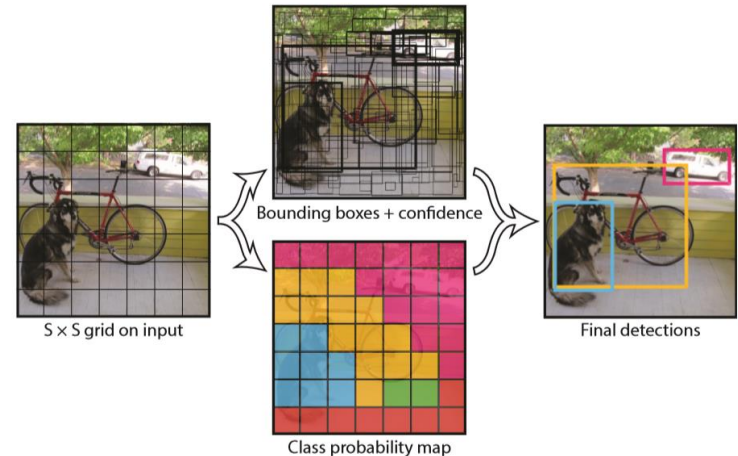
Fully Convolutional Network predicts  $B$  default bounding boxes per grid cell, with confidence.

Confidence is based on the Intersection over Union of the Bounding Boxes with each grid cell.

Closing a Non-Maxima Suppression is used to prune the detections to the final detections.

### You Only Look Once: Unified, Real-Time Object Detection

Joseph Redmon\*, Santosh Divvala\*<sup>†</sup>, Ross Girshick<sup>‡</sup>, Ali Farhadi\*<sup>†</sup>  
University of Washington\*, Allen Institute for AI<sup>†</sup>, Facebook AI Research<sup>‡</sup>  
<http://pjreddie.com/yolo/>



<https://arxiv.org/abs/1506.02640>

## SSD: Single Shot MultiBox Detector

Discretization of the input image into a  $S \times S$  grid.

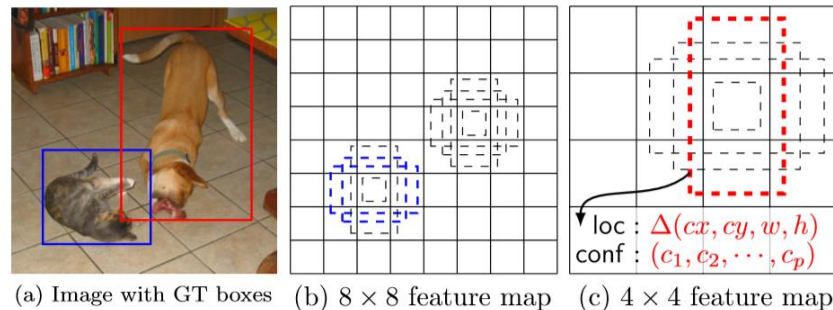
Fully Convolutional Network predicts B default bounding boxes per grid cell, with confidence.

Final detections are chosen with respect to their confidence

## SSD: Single Shot MultiBox Detector

Wei Liu<sup>1</sup>, Dragomir Anguelov<sup>2</sup>, Dumitru Erhan<sup>3</sup>, Christian Szegedy<sup>3</sup>,  
Scott Reed<sup>4</sup>, Cheng-Yang Fu<sup>1</sup>, Alexander C. Berg<sup>1</sup>

<sup>1</sup>UNC Chapel Hill <sup>2</sup>Zoox Inc. <sup>3</sup>Google Inc. <sup>4</sup>University of Michigan, Ann-Arbor  
wliu@cs.unc.edu, <sup>2</sup>drago@zoox.com, <sup>3</sup>{dumitru, szegedy}@google.com,  
<sup>4</sup>reedscot@umich.edu, <sup>1</sup>{cyfu, aberg}@cs.unc.edu



<https://arxiv.org/abs/1512.02325>

### RFCN: Region-based Fully Convolutional Network

Two stage object detection.

First stage is Fully Convolutional Network, such as ResNet-101 for region proposals

Second stage does classification on the max-pooled proposed regions

---

### R-FCN: Object Detection via Region-based Fully Convolutional Networks


---

Jifeng Dai  
Microsoft Research

Yi Li\*  
Tsinghua University

Kaiming He  
Microsoft Research

Jian Sun  
Microsoft Research



<https://arxiv.org/abs/1605.06409>

## TensorFlow Model Zoo

Collection of detection models pre-trained on different object detection datasets.



[https://github.com/tensorflow/models/blob/master/research/object\\_detection/g3doc/detection\\_model\\_zoo.md](https://github.com/tensorflow/models/blob/master/research/object_detection/g3doc/detection_model_zoo.md)

Model name	Speed (ms)	COCO mAP[^1]	Outputs
<a href="#">ssd_mobilenet_v1_coco</a>	30	21	Boxes
<a href="#">ssd_mobilenet_v1_0.75_depth_coco</a> ☆	26	18	Boxes
<a href="#">ssd_mobilenet_v1_quantized_coco</a> ☆	29	18	Boxes
<a href="#">ssd_mobilenet_v1_0.75_depth_quantized_coco</a> ☆	29	16	Boxes
<a href="#">ssd_mobilenet_v1_ppn_coco</a> ☆	26	20	Boxes
<a href="#">ssd_mobilenet_v1_fpn_coco</a> ☆	56	32	Boxes
<a href="#">ssd_resnet_50_fpn_coco</a> ☆	76	35	Boxes
<a href="#">ssd_mobilenet_v2_coco</a>	31	22	Boxes
<a href="#">ssd_mobilenet_v2_quantized_coco</a>	29	22	Boxes
<a href="#">ssdlite_mobilenet_v2_coco</a>	27	22	Boxes
<a href="#">ssd_inception_v2_coco</a>	42	24	Boxes
<a href="#">faster_rcnn_inception_v2_coco</a>	58	28	Boxes
<a href="#">faster_rcnn_resnet50_coco</a>	89	30	Boxes
<a href="#">faster_rcnn_resnet50_lowproposals_coco</a>	64		Boxes
<a href="#">rfcn_resnet101_coco</a>	92	30	Boxes
<a href="#">faster_rcnn_resnet101_coco</a>	106	32	Boxes
<a href="#">faster_rcnn_resnet101_lowproposals_coco</a>	82		Boxes
<a href="#">faster_rcnn_inception_resnet_v2_atrous_coco</a>	620	37	Boxes
<a href="#">faster_rcnn_inception_resnet_v2_atrous_lowproposals_coco</a>	241		Boxes
<a href="#">faster_rcnn_nas</a>	1833	43	Boxes
<a href="#">faster_rcnn_nas_lowproposals_coco</a>	540		Boxes

# Neural Network Object Detection - Datasets and benchmarking

## Common Objects in Context

COCO-Detection has 200k images with bounding boxes or pixel-wise labels

### Annotations


80 object categories (person, elephant, etc.), as well as captions.

### Number of samples

200000 bounding box level annotations

### Metric

Average Precision, AP across scales, Average Recall, AR across scales



<https://github.com/nightrome/cocostuff>



# Neural Network Object Detection - Datasets and benchmarking

## PASCAL Visual Object Classes

For each of twenty object classes predict the presence/absence of at least one object of that class in a test image.

### Annotations

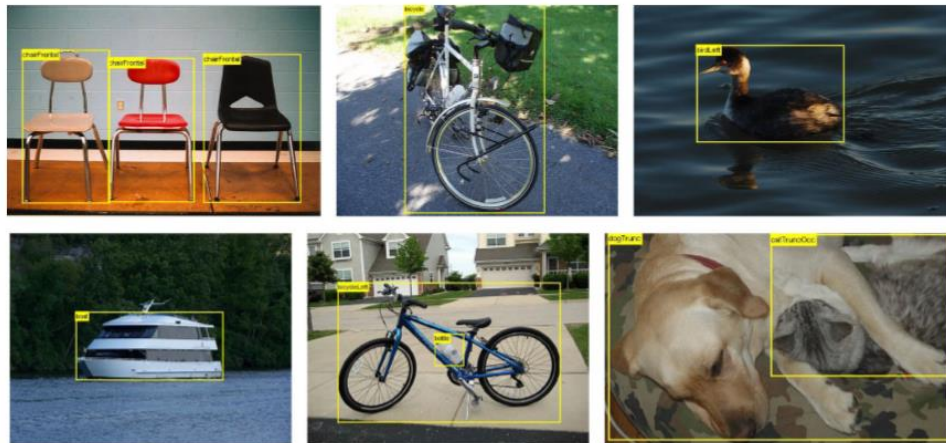
20 object classes (Person, Bicycle, etc.)

### Number of samples

11540 bounding box level annotations

### Metric

Average Precision, AP across scales,  
Average Recall, AR across scales



<http://host.robots.ox.ac.uk/pascal/VOC/pubs/everingham15.pdf>



# Neural Network Object Detection - Datasets and benchmarking

## KITTI

We take advantage of our autonomous driving platform Annieway to develop novel challenging real-world computer vision benchmarks.

### Annotations


2D bounding box annotations with classes

### Number of samples

7481 training images and  
7518 test images


### Metric

Average Precision,  
Average Recall, both across scales




<http://www.cvlibs.net/publications/Geiger2013IJRR.pdf>

### So how do we evaluate our Object Detectors?



### So how do we evaluate our Object Detectors?


Define when a detection is correct, a true positive and when it is not.  
A prediction will usually not overlap perfectly with the ground truth.



## So how do we evaluate our Object Detectors?

Define when a detection is correct, a true positive and when it is not.  
A prediction will usually not overlap perfectly with the ground truth.

Filter predictions based on  
the object detectors **confidence**



### So how do we evaluate our Object Detectors?

Define when a detection is correct, a true positive and when it is not.  
A prediction will usually not overlap perfectly with the ground truth.

Assign true positives based on  
**Intersection over Union** with the  
ground truth

$$IoU(A, B) = \frac{A \cap B}{A \cup B}$$

$$IoU \geq p, p \in (0, 1)$$

### So how do we evaluate our Object Detectors?

Define when a detection is correct, a true positive and when it is not.  
A prediction will usually not overlap perfectly with the ground truth.

Assign true positives based on  
**Intersection over Union**

```
prediction_prob = np.asarray([0.23, 0.37, 0.59, 0.84, 0.10, 0.92, 0.57, 0.49, 0.67, 0.51])
# Threshold for assignment
p = 0.5

# Define ground truth and true positives (1) and false negatives (0) of the prediction
groundtruth = np.asarray([1, 1, 1, 1, 1, 1, 1, 1, 0, 0])
prediction = (prediction_prob > p)*1
print(prediction)
# [0 0 1 1 0 1 1 0 1 1]
```

## Precision

Detections over all Predictions

True Positives over True Positives and False Positives

$$mAP = \frac{1}{|c|} \sum_c \frac{TP(c)}{TP(c) + FP(c)}$$

```
# -----  
# 1. Precision  
P = sum(prediction*groundtruth) / sum(prediction)  
print('Precision: ', P)  
# Precision: 0.66  
# -----
```

## Recall

Detections over the number of Object Instances.

True Positives over True Positives and False Negatives.

$$mAR = \frac{1}{|c|} \sum_c \frac{TP(c)}{TP(c) + FN(c)}$$

```
# -----  
# 2. Recall  
R = sum(prediction*groundtruth) / sum(groundtruth)  
print('Recall: ', R)  
# Recall: 0.5  
# -----
```



## F1-Score

Harmonic mean between recall and precision

Harmonic mean lays more emphasis on the weaker metric, than the arithmetic would.

$$F1 = 2 \frac{mAP \ mAR}{mAP + mAR}$$

```
# -----  
# 3. F1-Score  
# Harmonic mean  
F = 2*R*P / (R+P)  
print('F1-Score: ', F)  
# F1-Score: 0.57  
  
meanPR = (R+P) / 2  
print('Arithmetic mean: ', meanPR)  
# Arithmetic mean: 0.58  
# -----
```

### Confusion Matrix

The comparison of the model's predictions and the correct values in the form of a contingency table

```
# -----  
# 4. Confusion matrix  
CM = tf.confusion_matrix(labels=groundtruth,  
                          predictions=prediction)  
  
with tf.Session() as sess:  
    print('Confusion Matrix: ', sess.run(CM))  
# Confusion Matrix:  
# [ [0 2]  
#   [4 4] ]  
#  
# [ [true_negatives | false_positives]  
#   [false_negatives | true_positives ] ]  
# -----
```

## Confusion Matrix


Which number in the matrix does never apply, when dealing with Object Detectors?

```
# -----  
# 4. Confusion matrix  
CM = tf.confusion_matrix(labels=groundtruth,  
                          predictions=prediction)  
  
with tf.Session() as sess:  
    print('Confusion Matrix: ', sess.run(CM))  
# Confusion Matrix:  
# [ [0 2]  
#   [4 4] ]  
#  
# [ [true_negatives | false_positives]  
#   [false_negatives | true_positives ] ]  
# -----
```

## Receiver Operating Characteristic (ROC)

The ROC curve displays the trade off between true-positive-rate and true-negative-rate.


It can be seen as a matter of optimizing a threshold value with respect to a functional requirement.



## Precision Recall Curve (PRC)

The PR curve displays the trade off between Precision and Recall.

It can be seen as a matter of optimizing a threshold value with respect to a functional requirement.




## Receiver Operating Characteristic (ROC)

The ROC curve displays the trade off between Precision and Recall.

It can be seen as a matter of optimizing a **threshold** value with respect to a functional requirement.

- Intersection over Union
- Confidence



# References

- [2] Joseph Redmon, Santosh Kumar Divvala, Ross B. Girshick, and Ali Farhadi. You only look once: Unified, real-time object detection. *CoRR*, abs/1506.02640, 2015.
- [3] Wei Liu, Dragomir Anguelov, Dumitru Erhan, Christian Szegedy, Scott Reed, Cheng-Yang Fu, and Alexander C. Berg. Ssd: Single shot multibox detector. *14th european conference on computer vision*, pages 21–37, 2016.
- [4] Jifeng Dai, Yi Li, Kaiming He, and Jian Sun. R-FCN: object detection via region-based fully convolutional networks. *CoRR*, abs/1605.06409, 2016.
- [5] Tsung-Yi Lin, Michael Maire, Serge J. Belongie, Lubomir D. Bourdev, Ross B. Girshick, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C. Lawrence Zitnick. Microsoft COCO: common objects in context. *CoRR*, abs/1405.0312, 2014.
- [6] Mark Everingham, Luc Van Gool, Christopher K. I. Williams, John Winn, and Andrew Zisserman. The pascal visual object classes (voc) challenge. *International Journal of Computer Vision*, 88(2):303–338, Jun 2010.
- [7] Andreas Geiger, Philip Lenz, and Raquel Urtasun. Are we ready for autonomous driving? the kitti vision benchmark suite. In *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2012.
- [8] Jesse Davis and Mark Goadrich. The relationship between precision-recall and roc curves. In *Proceedings of the 23rd International Conference on Machine Learning, ICML '06*, pages 233–240, New York, NY, USA, 2006. ACM.
- [9] Mark Everingham, Luc Van Gool, Christopher K. I. Williams, John Winn, and Andrew Zisserman. The pascal visual object classes (voc) challenge. *International Journal of Computer Vision*, 88(2):303–338, Jun 2010.

Object Detection with neural networks

## **Introduction to Transfer Learning**

Transfer Learning a problem statement


Settings for Transfer Learning

Transfer Learning with neural networks



## Introduction – Transfer Learning a problem statement


*In machine learning it is assumed that the **training** data and the data received during **inference** are drawn from the same distribution and domain.*



## Introduction – Transfer Learning a problem statement

*In machine learning it is assumed that the training data and the data received during inference are drawn from the same distribution and domain.*


*When facing real-world applications this assumption is regularly violated (e.g. day and night, new car models, etc.)*



## Introduction – Transfer Learning a problem statement


*When facing real-world applications this assumption is regularly violated (e.g. day and night, new car models, etc.)*

*We distinguish between source and target domain.*



## Introduction – Transfer Learning a problem statement


**Definition 1.1 (Domain).** A domain  $D$  consists of a feature space  $\mathcal{X}$  and a marginal probability distribution  $P(X)$ , where  $X = \{x_1, \dots, x_n\} \in \mathcal{X}$ . The source domain data is denoted as  $D_S = \{(x_{S_1}, y_{S_1}), \dots, (x_{S_n}, y_{S_n})\}$ . The target domain data is denoted as  $D_T = \{(x_{T_1}, y_{T_1}), \dots, (x_{T_n}, y_{T_n})\}$  [19].



## Introduction – Transfer Learning a problem statement


*Machine learning models are trained with respect to a task  $T$  such as object detection, segmentation, etc.*

*These can also vary between source and target.*



## Introduction – Transfer Learning a problem statement

**Definition 1.2 (Task).** A task  $T$  consists of a label space  $\mathcal{Y}$  and an objective prediction function  $f(\cdot)$ , denoted by  $T = \{\mathcal{Y}, f(\cdot)\}$ .  $f(x)$  can be written as  $P(y|x)$  [19].



## Introduction – Transfer Learning a problem statement

*Transfer Learning is to be understood with respect to the concerned domains and their combination with a task setting*

## Introduction – Transfer Learning a problem statement

*Transfer Learning is to be understood with respect to the concerned domains and their combination with a task setting*

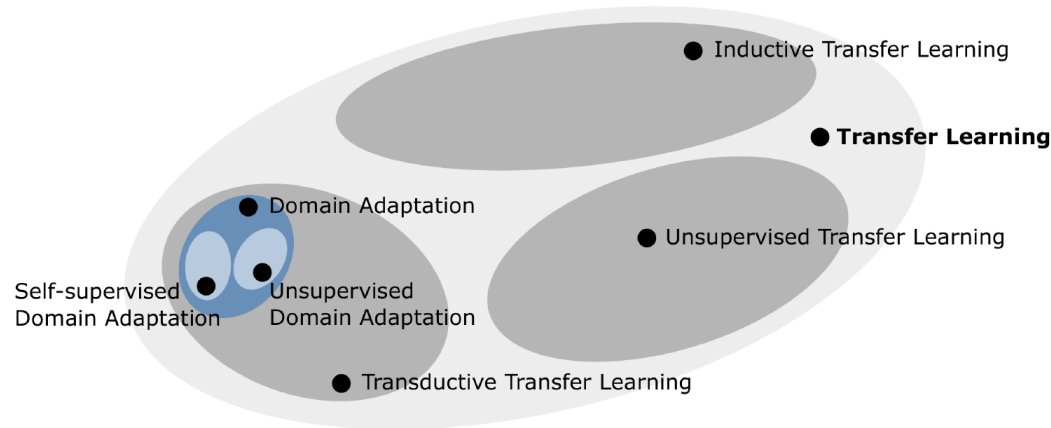


Figure 1.5: Overview of the different transfer learning settings: Inductive transfer learning, unsupervised transfer learning, and transductive transfer learning.



## Introduction – Transfer Learning a problem statement

*Transfer Learning is to be understood with respect to the concerned domains and their combination with a task setting*

- **Inductive** Transfer Learning
- **Unsupervised** Transfer Learning
- **Transductive** Transfer Learning

## Introduction – Transfer Learning a problem statement

*Transfer Learning is to be understood with respect to the concerned domains and their combination with a task setting*

### **Inductive Transfer Learning**

Labels are available in the source and target domain.

The target task is different from the source task. Can be understood as task adaptation.

Unsupervised Transfer Learning

Transductive Transfer Learning

## Introduction – Transfer Learning a problem statement

*Transfer Learning is to be understood with respect to the concerned domains and their combination with a task setting*

Inductive Transfer Learning

### **Unsupervised Transfer Learning**

There are no labels available. And the tasks in source and target are different.  
Think clustering approaches.

Transductive Transfer Learning

## Introduction – Transfer Learning a problem statement

*Transfer Learning is to be understood with respect to the concerned domains and their combination with a task setting*

Inductive Transfer Learning

Unsupervised Transfer Learning

### **Transductive Transfer Learning**

There are only labels in the source domain and none in the target domain. The tasks are the same in both domains. Can be understood as domain adaptation.

# References

- [1] S. J. Pan and Q. Yang. A survey on transfer learning. *IEEE Transactions on Knowledge and Data Engineering*, 22(10):1345–1359, Oct 2010.

Object Detection with neural networks  
Introduction to Transfer Learning

**Transfer Learning with neural networks**

Fine Tuning

Freezing Weights

Shared Latent Feature Spaces

Pseudo Labels

Overview state-of-the-art approaches

Datasets and Benchmarking

Deep dive Day2Night

## Transfer Learning with Neural Networks

The intuition behind transfer learning is that if a model trained on a large and general enough dataset, this model will effectively serve as a **generic model of the visual world**.

## Transfer Learning with Neural Networks

The intuition behind transfer learning is that if a model trained on a large and general enough dataset, this model will effectively serve as a generic model of the visual world.

*You can then take advantage of these learned feature maps without having to start from scratch training a large model on a large dataset.*



## Transfer Learning with Neural Networks - Fine Tuning

A **Pre-trained model** is a model that is trained on the source domain for a source task.



Pre-trained  
model

Car

## Transfer Learning with Neural Networks - Fine Tuning

A **Pre-trained model** is a model that is trained on source domain for a source task.

Think **Object Detection Model Zoo**



Pre-trained  
model

Car

[https://github.com/tensorflow/models/blob/master/research/object\\_detection/g3doc/detection\\_model\\_zoo.md](https://github.com/tensorflow/models/blob/master/research/object_detection/g3doc/detection_model_zoo.md)

## Transfer Learning with Neural Networks - Fine Tuning

A **Pre-trained model** is a model that is trained on the source domain for a source task.

The source domain is usually a **large dataset**.



Pre-trained  
model

Car

## Transfer Learning with Neural Networks - Fine Tuning

A **Pre-trained model** is a model that is trained on the source domain for a source task.

The source domain is usually a **large dataset**.

Think **COCO**, **ImageNet**, **PASCAL VOC**



Pre-trained  
model

Car

## Transfer Learning with Neural Networks - Fine Tuning

You either use a **pre-trained model as it is**.



Pre-trained  
model

Car

## Transfer Learning with Neural Networks - Fine Tuning

You either use a pre-trained model as it is.

- No additional training
- Usually **performance drops**




Pre-trained  
model

Car

## Transfer Learning with Neural Networks - Fine Tuning


You can **fine-tune** a **pre-trained model** by retraining on additional data.



## Transfer Learning with Neural Networks - Fine Tuning

You can fine-tune a pre-trained model by retraining on additional data.

- Additional work to set up retraining pipeline
- Usually performance drops in the source domain
- **Performance improvement** in the target domain





## Transfer Learning with Neural Networks - Feature Extraction

You can **repurpose** the **pre-trained layers** as feature extraction layers for low level features.

And adding layers that learn the required high level features on the new data.



Pre-trained  
layers

Newly-trained  
layers


Car

# Transfer Learning with Neural Networks - Feature Extraction

You can repurpose the pre-trained layers as feature extraction layers for low level features.

And adding layers that learn the required high level features on the new data.

- Additional work to assemble layer structure
- **Regularization** effects
- **Performance improvement** in the target domain
- **Output** layer becomes **adjustable** (such as for adding a class)



## Transfer Learning with Neural Networks - Shared Latent Feature Space

So far the transfer was **supervised**, meaning we were in possession of labels to deploy in the target domain.

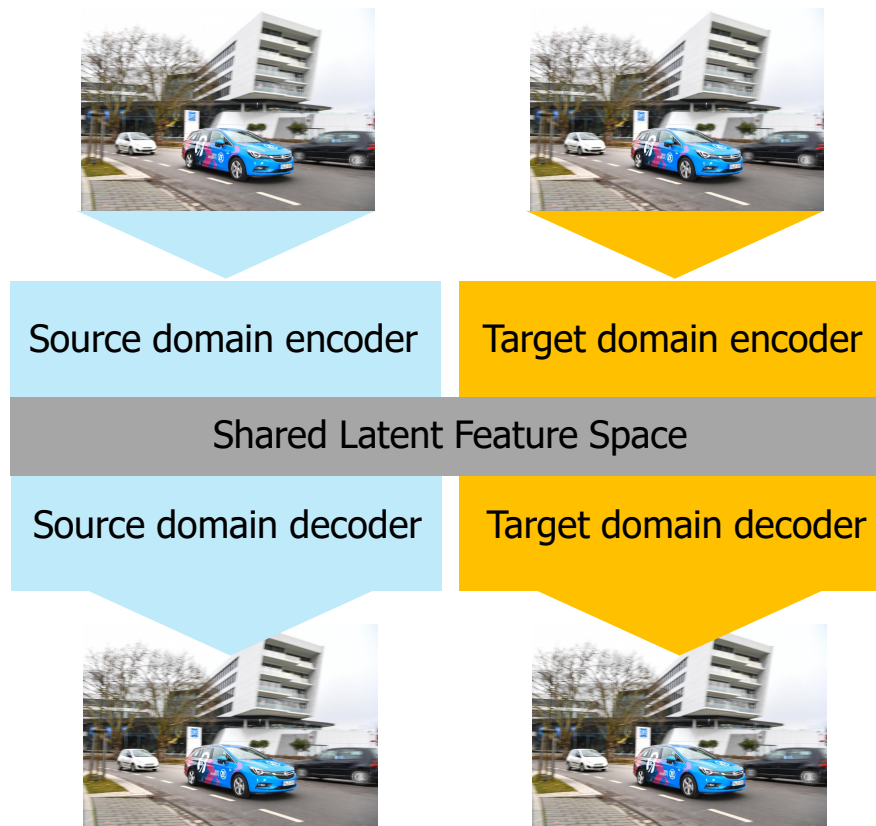
## Transfer Learning with Neural Networks - Shared Latent Feature Space

So far the transfer was supervised, meaning we were in possession of labels to deploy in the target domain.

Next up are methods that do not require target domain ground truth – **unsupervised transfer learning** approaches.

## Transfer Learning with Neural Networks - Shared Latent Feature Space

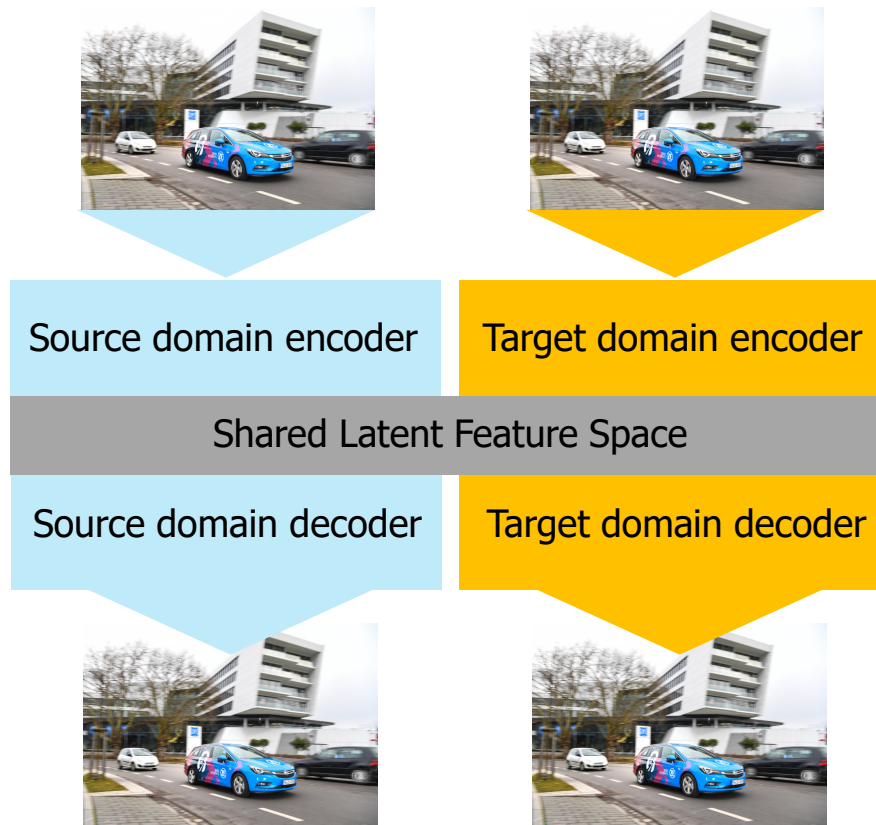
We can introduce a **shared layer** which models a shared latent feature space within an encoder decoder architecture.



# Transfer Learning with Neural Networks - Shared Latent Feature Space

We can introduce a shared layer which models a shared latent feature space within an encoder decoder architecture.

This layer **shares** the **parameters across** the two architectures



## Transfer Learning with Neural Networks - Shared Latent Feature Space

We can introduce a shared layer which models a shared latent feature space within an encoder decoder architecture.

This layer **shares** the **parameters across** the two architectures and thus develops a shared **feature representation** for both domains.



Source domain encoder



Target domain encoder

Shared Latent Feature Space

## Transfer Learning with Neural Networks - Shared Latent Feature Space

We can introduce a shared layer which models a shared latent feature space within an encoder decoder architecture.

The shared **feature representation** is used as a base for training the task on source domain labels.



Source domain encoder

Shared Latent Feature Space

Layers trained on Source labels

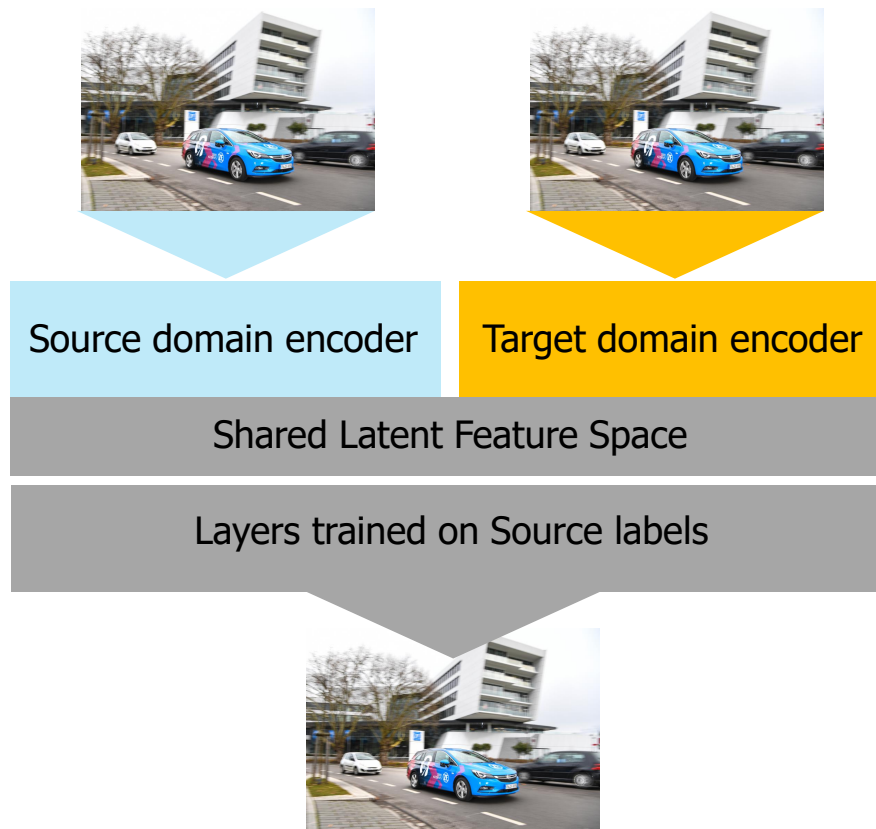




## Transfer Learning with Neural Networks - Shared Latent Feature Space

We can introduce a shared layer which models a shared latent feature space within an encoder decoder architecture.

During **inference** the prediction is done **independent** of the sample's **domain**.



## Transfer Learning with Neural Networks - Self-supervised

In case the target domain is **overlapping** with the source domain or can be brought in a **continuous composition** the model can transfer the knowledge **self-supervised**.



Pre-trained  
model

Car

## Transfer Learning with Neural Networks - Self-supervised

In case the target domain is overlapping with the source domain or can be brought in a continuous composition the model can transfer the knowledge self-supervised.

This is realized by introducing **pseudo labels** on the target domain data, which are predicted by the pre-trained model.

$$y_{P,i} = \begin{cases} 1, & \text{if } i = \arg \max_i \theta(x_T). \\ 0, & \text{else.} \end{cases}$$



Pre-trained  
model

Pseudo Label

## Transfer Learning with Neural Networks - Self-supervised

In case the target domain is overlapping with the source domain or can be brought in a continuous composition the model can transfer the knowledge self-supervised.

This is realized by introducing **pseudo labels** on the target domain data, which are predicted by the pre-trained model.

$$y_{P,i} = \begin{cases} 1, & \text{if } i = \arg \max_i \theta(x_T). \\ 0, & \text{else.} \end{cases}$$



Pre-trained  
model

Pseudo Label



Re-trained  
model

Car

## Semi-Supervised Self-Training of Object Detection Models

### Semi-supervised self-training of object detection models

Train a model on the source domain (supervised).

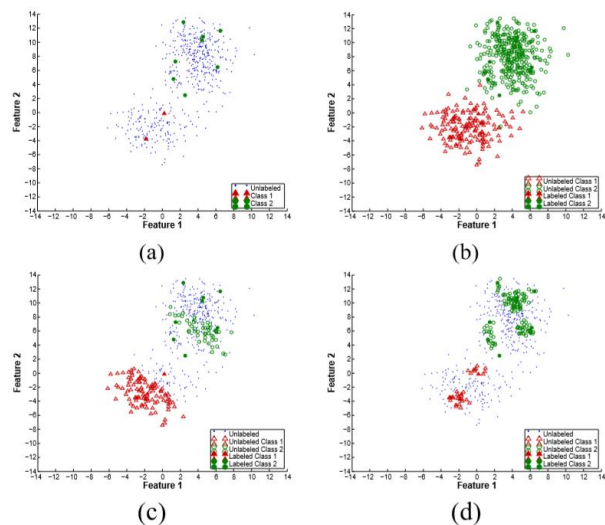
Generate Pseudo-Labels on target domain.

Add samples to the training set for fine-tuning with respect to a prediction confidence threshold.

Chuck Rosenberg  
Google, Inc.  
Mountain View, CA 94043  
chuck@google.com

Martial Hebert  
Carnegie Mellon University  
Pittsburgh, PA 15213  
hebert@ri.cmu.edu

Henry Schneiderman  
Carnegie Mellon University  
Pittsburgh, PA 15213  
hws@ri.cmu.edu



[https://www.ri.cmu.edu/pub\\_files/pub4/rosenberg\\_charles\\_2005\\_1/rosenberg\\_charles\\_2005\\_1.pdf](https://www.ri.cmu.edu/pub_files/pub4/rosenberg_charles_2005_1/rosenberg_charles_2005_1.pdf)

## Learning without Forgetting


Train a model on a source task (supervised).

Generate pseudo-labels for the source task on the target domain.

Add nodes for target task and jointly train for both tasks.

## Learning without Forgetting

Zhizhong Li, Derek Hoiem, *Member, IEEE*



<https://arxiv.org/pdf/1606.09282.pdf>

## Unsupervised image-to-image translation network UNIT


Map samples from different domains (source and target) into a shared feature space.

The mapping is learned by a joint encoding and decoding as well as an adversarial loss of two discriminators.

Later samples from one domain can be translated into the other domain and respectively.

## Unsupervised Image-to-Image Translation Networks

Ming-Yu Liu, Thomas Breuel, Jan Kautz  
NVIDIA  
{mingyul,tbreuel,jkautz}@nvidia.com



<https://arxiv.org/pdf/1703.00848.pdf>

### Model-specific

- Task-specific performance metrics are considered for source domain, target domain, and intermediate domains
- Catastrophic forgetting depicting the performance decrease on the source domain.
- Domain gap depicting the performance difference between source and target domain.

### Training-specific



Model-specific

## Training-specific

- Model size efficiency as in memory size increase of the model
- Sample size efficiency as in memory size when samples need to be stored for retraining and replay strategies
- Computational efficiency as in operations needed to adapt a model to the next domain

# References

- [10] Chuck Rosenberg, Martial Hebert, and Henry Schneiderman. Semi-supervised self-training of object detection models. In *2005 Seventh IEEE Workshops on Applications of Computer Vision (WACV/MOTION'05)-Volume 1*, volume 1, pages 29–36. IEEE, 2005.
- [11] Z. Li and D. Hoiem. Learning without forgetting. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(12):2935–2947, Dec 2018.
- [12] Ming-Yu Liu, Thomas Breuel, and Jan Kautz. Unsupervised image-to-image translation networks. In I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances in Neural Information Processing Systems 30*, pages 700–708. Curran Associates, Inc., 2017.
- [13] Natalia Díaz-Rodríguez, Vincenzo Lomonaco, David Filliat, and Davide Maltoni. Don't forget, there is more than forgetting: new metrics for continual learning. *arXiv preprint arXiv:1810.13166*, 2018.
- [14] Sebastian Farquhar and Yarin Gal. Towards robust evaluations of continual learning. *arXiv preprint arXiv:1805.09733*, 2018.

## Tutorial

jupyter TL-TUT Last Checkpoint: 08/13/2019 (autosaved)

Logout

File Edit View Insert Cell Kernel Help Trusted tensorflow

Applied Deep Learning Tutorial

Transfer Learning for Object Classification

Introduction

In this tutorial, you will attempt to benefit from a model that has been pretrained for the same task but on a different dataset. You will deploy the first layers and their feature extraction capabilities of a converged network. This process is known as transfer learning.




Fig. 1: Cat and dog in an image

Core idea

A pre-trained model is a saved network that was previously trained on a large dataset, typically on a large-scale image-classification task, such as [ImageNet](#), and [COCO](#). We can either use the pretrained model as it is for inference on the task it has been trained on or we can do transfer learning using the pretrained convnets for further training on a new dataset with a possibly new output space.

The intuition behind transfer learning is that if this model trained on a large and general enough dataset, this model will effectively serve as a generic model of the visual world and the semantic features present in the visual world and shared between all visual tasks. We can leverage these learned feature maps without having to train a large model on a large dataset by using these models as the basis of our own model specific to our task. There are 2 scenarios of transfer learning using a pretrained model.

- Fine Tuning or Retraining: Unfreezing a few of the top layers of a frozen model base used for feature extraction, and jointly training both the newly added classifier layers as well as the last layers of the frozen model. This allows us to "fine tune" the higher order feature representations in addition to our final classifier in order to make them more relevant for the specific task involved.
- Feature Extraction: Use the representations learned by a previous model to extract meaningful features from new samples. We simply add a new output layer, which will be trained from scratch, on top of the pretrained model so that we can repurpose the feature maps learned previously for our dataset and our new output space.

# Thanks for your time

## Questions?

**Contact! [Subject]**

mark.schutera@kit.edu [ADL]

