

Applied Deep Learning Segmentation Networks

Mark Schutera

Karlsruhe Institute of Technology | ZF Friedrichshafen AG | Hochschule Ravensburg-Weingarten

1. Deep Learning Foundations
-
3. Transfer Learning and Object Detection
-
- 5. Segmentation Networks**
-
8. Deep Reinforcement Learning
-
10. Generative Adversarial Networks
-
12. Recurrent Neural Networks

Intermediate presentation

Ten minutes on achievements, problems, next steps

Due 7. Lecture

Final presentation

Code and results

Due 14. Lecture

Final documentation

Paper and code on github or jupyter notebook

Deadline 14. Lecture

Course features

Sli.do

Every question matters.

Get the app.

Ask questions (with slide number)
or vote on other students' questions
during the lecture.

And give direct feedback.

#TOBEDETERMINED

Questions will be covered
immediately or in the next lecture in
more depth.

Github

Find slides, tutorials, flashcards and
references on Github.

<https://github.com/schutera/DeepLearningLecture> **Schutera**

You found typos, additional
material such as links, algorithms,
papers, literature or want to
contribute to the slides and lecture
notes..

..Feel free to contribute, e-mail me.

Course features

Sli.do

Every question matters.

Get the app.

Ask questions (with slide number)
or vote on other students' questions
during the lecture.

#TOBEDETERMINED

Questions will be covered directly or
in the next lecture in more depth.

Github

Typos, additional material such as
links, algorithms, paper, literature,
lecture notes..

..Feel free to contribute.

Grade Bonus .3

Prepare flashcards based on Ian
Goodfellow's Deep Learning Book

- Commit to flashcard set by
emailing me, first come first serve
- Must be comprehensive

Introduction to segmentation

Segmentation a problem statement

Applications for segmentation

Conventional segmentation approaches


Segmentation with neural networks

Introduction – Segmentation a problem statement


Classification



Classification
+ Localization



Object Detection



Object Tracking



Introduction – Segmentation a problem statement

Segmentation is classification on pixel-level, which results in super-pixels or segments or groups of pixels based on some criteria.

Semantic Segmentation



Instance Segmentation



Autonomous Driving

- Scene understanding
- Understanding of shapes
- Supports sensor fusion with point cloud sensors
- Free space detection

Geo Analytics

Medical Imaging




<https://www.cityscapes-dataset.com/>

Introduction – Applications for segmentation

Autonomous Driving

Geo Analytics

- Building structures
- Road network analysis
- Wildfire detection
- Water supply tracking
- Real time crisis management
- Weather prediction



<https://github.com/mapbox/robosat>


Medical Imaging

Introduction – Applications for segmentation

Autonomous Driving
Geo Analytics

Medical Imaging

- Tissue localization and analysis
- Volume approximations
- Surgery planning
- Temporal tumor or tissue development
- Tooling for drug testing



<https://osf.io/snb6p/>

Introduction – Conventional segmentation approaches

Image segmentation is a well researched field.

In order to design neural networks it is a good thing to really understand the task at hand.

Thresholding

Edge detection

Clustering

Region growing

Introduction – Conventional segmentation approaches

Thresholding

The simplest method of image segmentation is called the thresholding method.


This method is based on a threshold value to turn a gray-scale image into a binary image (mask).

Usually this is just one step of many.

Edge detection

Clustering

Region growing



Thresholding on Zebrafish for eye segmentation

Lena Test image

Lena , the 'hello world!' of image processing. 330x330

Cover photo of 1972 Playboy magazine of the Swedish model Lena Söderberg.

Since then Lena was a guest at several IEEE conferences. The image also sparked discussions on gender-equality in the male-dominated field of engineering.

It is a good test image because of its detail, flat regions, shading, and texture.



<https://en.wikipedia.org/wiki/Lenna>

Introduction – Conventional segmentation approaches

Thresholding

Edge detection

Segment boundaries and edges are closely related.

Since there is often a large gradient at the segment boundaries.



Canny Edge Detection

$$L_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} L \quad \text{and} \quad L_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} L.$$

Sobel Operators for Edge Detection

Clustering

Region growing

Introduction – Conventional segmentation approaches

Thresholding

Edge detection

Often edge detectors are combined with morphological operators to close the detected edges.

Clustering

Region growing



Edge Detection for segmentation

Introduction – Conventional segmentation approaches

Thresholding

Edge detection


Clustering (Color quantization)

K-means with 3 features (R,G,B) and K centroids.

The centroids are iteratively adjusted until convergence.

After the clustering, the centroid values are applied to the pixels in their cluster.

Region growing



Clustering for K=4 (top) and K=8 (bottom)

Introduction – Conventional segmentation approaches

Thresholding
Edge detection
Clustering



Thresholding, Find valleys, Region growing for Segmentation

Region growing

Any image can be viewed as a topographic surface due to the gradients in the image. You start by finding valleys and filling them with different colored water (labels). As the water rises, depending on the peaks (gradients) nearby, water from different valleys, will get in contact, that is where you build barriers. The barriers give the resulting segmentation

References

- [1] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [2] Andrej Karpathy. Cs231n: Convolutional neural networks for visual recognition. <http://cs231n.github.io/neural-networks-3/>, 2018. Zugriff: 20.01.2018.
- [3] Schutera, Mark, Steffen Just, Jakob Gierden, Ralf Mikut, Markus Reischl, and Christian Pylatiuk. 2019. "Machine Learning Methods for Automated Quantification of Ventricular Dimensions." OSF. March 28. osf.io/snb6p.
- [4] Mark Schutera, Thomas Dickmeis, Marina Mione, Ravindra Peravali, Daniel Marcato, Markus Reischl, Ralf Mikut, and Christian Pylatiuk. Automated phenotype pattern recognition of zebrafish for high-throughput screening. *Bioengineered*, 7(4):261–265, 2016.
- [5] Canny, J., *A Computational Approach To Edge Detection*, IEEE Transactions on Pattern Analysis and Machine Intelligence, 8(6):679–698, 1986.

Introduction to segmentation

Segmentation with neural networks

Basic structure

Overview state-of-the-art

Datasets and benchmarking

Deep dive U-Net

Why classical segmentation approaches

- Interpretability
- Only a few samples needed
- No labeling needed
- No training needed
- Usually better runtime during inference

Why not?



<https://osf.io/snb6p/>

Why classical segmentation approaches

Why segmentation by neural networks?

- Do generalize better
- Feature engineering has a limited capacity to capture semantics
- Feature engineering is expensive and time consuming



<https://osf.io/snb6p/>

Feature Representation by Convolution

Idea is to classify each pixel of an input image by representation learning

Downsampling

Upsampling

Parameter sharing

Feature Representation by Convolution


A convolutional layer is not fully connected, but has a narrowed down receptive field (e.g. 3×3).

The parameters of each filter are spatially shared: A feature that is useful in one place, ought to be useful in another, too.

Downsampling

Upsampling

Parameter sharing



- **Depth:** number of filters
- **Stride:** filter step size (when we “slide” it)
- **Padding:** zero-pad the input

<https://selfdrivingcars.mit.edu/>

Feature Representation by Convolution

0	0	0	0	0	0	0
0	1	2	0	0	0	0
0	2	2	0	0	0	0
0	2	2	2	1	0	0
0	1	0	0	2	0	0
0	1	1	0	1	0	0
0	0	0	0	0	0	0

Zero-Padding is adding zero-valued pixel to the image border (gray area).

Downsampling

Upsampling

Parameter sharing

Feature Representation by Convolution

0	0	0	0	0	0	0
0	1	2	0	0	0	0
0	2	2	0	0	0	0
0	2	2	2	1	0	0
0	1	0	0	2	0	0
0	1	1	0	1	0	0
0	0	0	0	0	0	0

Zero-padded image

0	0	-1
-1	0	0
-1	-1	-1

Filter

0

Bias

-4	-4	0
-3	-4	-3
0	-3	-1

Output

Downsampling

Upsampling

Parameter sharing

Feature Representation by Convolution

0	0	0	0	0	0	0
0	1	2	0	0	0	0
0	2	2	0	0	0	0
0	2	2	2	1	0	0
0	1	0	0	2	0	0
0	1	1	0	1	0	0
0	0	0	0	0	0	0

0	0	-1
-1	0	0
-1	-1	-1

Filter

0

Bias

-4	-4	0
-3	-4	-3
0	-3	-1

Output

Amount of filters or convolution depth: 1

Filter step size or Stride: 2

Zero-padded image

Downsampling

Upsampling

Parameter sharing

Feature Representation by Convolution

0	0	0	0	0	0	0
0	1	2	0	0	0	0
0	2	2	0	0	0	0
0	2	2	2	1	0	0
0	1	0	0	2	0	0
0	1	1	0	1	0	0
0	0	0	0	0	0	0

0	0	-1
-1	0	0
-1	-1	-1

Filter

0

Bias

-4	-4	0
-3	-4	-3
0	-3	-1

Output

Review edge detector:

Similar idea, now the parameters of the filters are learned. We want a lot of filters!

Downsampling

Upsampling

Parameter sharing

$$L_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} L \quad \text{and} \quad L_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} L.$$

Feature Representation by Convolution

0	0	0	0	0	0	0
0	1	2	0	0	0	0
0	2	2	0	0	0	0
0	2	2	2	1	0	0
0	1	0	0	2	0	0
0	1	1	0	1	0	0
0	0	0	0	0	0	0

Downsampling

Upsampling

Parameter sharing

0	0	-1
-1	0	0
-1	-1	-1

Filter

0


Bias

-4	-4	0
-3	-4	-3
0	-3	-1

Output

Review edge detector:

Similar idea, now the parameters of the filters are learned. And we want to go deep!



Convolutions

Downsampling

Convolutions at original image resolution are computational expensive:

Filter dimensions \times image dimensions \times number of filters \times number of input channels.

Motivating a convolutional
encoder-decoder
structure and Downsampling.

Upsampling

Parameter sharing


Neural Network Segmentation - Basic structure

Convolutions

Downsampling

*Convolutions at original image resolution are computational expensive:
Filter dimensions \times image dimensions \times number of filters \times number of input channels.*

Motivating a convolutional
encoder-decoder
structure and Downsampling.



<http://cs231n.github.io/>

Upsampling


Parameter sharing

Neural Network Segmentation - Basic structure


Convolutions

Downsampling

Strided convolutions



Filter 3x3x1



Zero-padded image

Upsampling


Parameter sharing

Neural Network Segmentation - Basic structure


Convolutions

Downsampling

Strided convolutions




Filter 3x3x1



Zero-padded image

Stride 1



Upsampling


Parameter sharing

Neural Network Segmentation - Basic structure


Convolutions

Downsampling

Strided convolutions




Filter 3x3x1




Zero-padded image

Stride 1



Stride 2




Upsampling

Parameter sharing


Neural Network Segmentation - Basic structure

Convolutions

Downsampling Strided convolutions




Filter 3x3x1




Zero-padded image

Stride 1



Stride 2



Stride 4




Upsampling

Parameter sharing

Neural Network Segmentation - Basic structure


Convolutions

Downsampling Max Pooling



Max Pooling
3x3 Stride 3

5	4	6	3	1	7
3	4	1	6	4	4
6	2	5	6	6	4
2	6	9	8	6	3
4	8	5	6	8	2
3	1	7	8	6	3



6	7
9	8

Upsampling

Parameter sharing


Neural Network Segmentation - Basic structure

Convolutions

Downsampling Max Pooling

Intuition is to decrease the resolution while keeping the strongest features of each channel.

Introducing a location invariance.



<https://selfdrivingcars.mit.edu/>

Upsampling

Parameter sharing

Neural Network Segmentation - Basic structure

Convolutions

Downsampling

Upsampling

Classification needs to happen in original image resolution


Motivating Upsampling inside the network structure.

Parameter sharing



Neural Network Segmentation - Basic structure

Convolutions
Downsampling

Upsampling Nearest neighbor



3x3 Stride 3




Parameter sharing

Neural Network Segmentation - Basic structure


Convolutions
Downsampling

Upsampling Bed of Nails



3x3 Stride 3

6	7
9	8



6	0	0	7	0	0
0	0	0	0	0	0
0	0	0	0	0	0
9	0	0	8	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Parameter sharing

Neural Network Segmentation - Basic structure

Convolutions
Downsampling


Upsampling Max Unpooling

Corresponding pairs of
downsampling and
upsampling layers.

Use position of pooling
layer for unpooling

Parameter sharing

5	4	3	3	1	7
3	4	1	6	4	4
6	2	5	6	6	4
2	6	9	7	6	3
4	8	5	6	8	2
3	1	7	7	6	3




6	7
9	8

...

...

4	5
7	2



0	0	0	0	0	5
0	0	0	0	0	0
4	0	0	0	0	0
0	0	7	0	0	0
0	0	0	0	2	0
0	0	0	0	0	0

Neural Network Segmentation - Basic structure

Convolutions

Downsampling





Upsampling

Transpose convolution

Learnable Upsampling, also known as: Upconvolution, or Deconvolution (bad terminology)

Stride: 3

Padding: 1



4x4

Parameter sharing

Neural Network Segmentation - Basic structure

Convolutions

Downsampling

Upsampling

Skip connections

Trade-off between classification and localization

- High level features from later in the network, enable high classification performance, since they are more discriminative and contain more useful semantic information.
- On the other hand, those deep features have low resolution and, thus pose a problem for localization performance.

Neural Network Segmentation - Basic structure

Convolutions


Downsampling

Upsampling

Skip connections

Combining low-level features,
which have high localization
accuracy

With the high-level features,
which have are descriptive but
low-resolution.




<http://cs231n.github.io/>

Last layer

Last Layer results in a tensor with $H \times W$ image resolution and a depth of C : Number of classes to segment.

The last layer should encode the values into a range of values of $(0;1)$.
Either by **softmax** or **sigmoid** function.



<http://cs231n.github.io/>


Cross-entropy
Dice-coefficient

Sigmoid

$$f(z) = \frac{1}{1 + e^{-z}}$$
$$f'(z) = f(z)(1 - f(z))$$

Binary classification only.

The probability sum does not need to be one.



<http://cs231n.github.io/neural-networks-1/>

Neural Network Activation Functions - Review

Softmax

Normalized exponential function

$$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

Used for multi-class segmentation.

Probability sum will be 1.

before after

2.0	0.7
1.0	0.2
0.1	0.1

Last layer

Binary Cross-entropy

$$J = - \frac{1}{N} \sum_{n=1}^N \left[y_n \log \hat{y}_n + (1 - y_n) \log(1 - \hat{y}_n) \right]$$

Heavily penalizes opposing predictions.

And gives rise to the problem of imbalanced classes.

Dice-coefficient

Neural Network Segmentation - Optimization


Last layer

Binary Cross-entropy


$$J = - \frac{1}{N} \sum_{n=1}^N \left[y_n \log \hat{y}_n + (1 - y_n) \log(1 - \hat{y}_n) \right]$$

Heavily penalizes opposing predictions.


And gives rise to the problem of imbalanced classes.



prediction



ground
truth



error 0.33

Dice-coefficient

Last layer

Binary Cross-entropy

$$dice\ loss = 1 - \frac{2 \sum(\tilde{\mathbf{y}} \odot \mathbf{y}_{seg}) + 1}{\sum(\tilde{\mathbf{y}}^2) + \sum(\mathbf{y}_{seg}^2) + 1}.$$

Dice-coefficient

Similar to the IoU (Intersection over union), and thus easy to interpret.

+1 is a smoothing factor for numeric stability.

More robust with respect to imbalanced classes.

Neural Network Segmentation - Optimization

Last layer


Binary Cross-entropy

Dice-coefficient

Similar to the IoU (Intersection over union), and thus better interpretability.

More robust with respect to imbalanced classes.

$$dice\ loss = 1 - \frac{2 \sum(\tilde{\mathbf{y}} \odot \mathbf{y}_{seg}) + 1}{\sum(\tilde{\mathbf{y}}^2) + \sum(\mathbf{y}_{seg}^2) + 1}.$$




Thought experiment

Assumption:

The maximum number of class 1 pixels in a single sample is 1.

This simulates an extreme class imbalance ratio of 1 to 8.




Thought experiment


Assumption:

The maximum number of class 1 pixels in a single sample is 1

What is the expected error if the model tries to predict a single pixel of class 1 and fails?



prediction




ground
truth

Thought experiment


Assumption:

The maximum number of class 1 pixels in a single sample is 1

What is the expected error if the model tries to predict a single pixel of class 1 and fails?



prediction



ground
truth

BCE
0.22

DL
0.66


Thought experiment

Assumption:


The maximum number of class 1 pixels in a single sample is 1

What is the expected error if the model tries to predict a single class 1?

What is the maximal expected error if the model predicts class 2 only?



prediction



ground
truth


Thought experiment

Assumption:


The maximum number of class 1 pixels in a single sample is 1

What is the expected error if the model tries to predict a single class 1?

What is the maximal expected error if the model predicts class 2 only?



prediction



ground
truth

BCE
0.11

DL
0.5

Thought experiment


How high is the pressure to get locked in a local minimum if predictions are initially random?

Thought experiment


How high is the pressure to get locked in a local minimum if predictions are initially random?

Binary cross-entropy

$$\begin{aligned} 0.22 * 8 + 0.00 * 1 &= 1.76 \\ 0.11 * 9 &= 1.00 \end{aligned}$$



43 %




Thought experiment


How high is the pressure to get locked in a local minimum if predictions are initially random?

Dice-loss

$$\begin{aligned} 0.66 * 8 + 0.00 * 1 &= 5.28 \\ 0.50 * 9 &= 4.50 \end{aligned}$$



15 %



How to deal with imbalanced classes

Choose **dice-loss** over cross-entropy.

How to deal with imbalanced classes

Choose dice-loss over cross-entropy.

Balance your cross-entropy according to the class imbalance.

In our case $\beta = 7/8$

$$\text{BCE}(p, \hat{p}) = -(\beta p \log(\hat{p}) + (1 - \beta)(1 - p) \log(1 - \hat{p}))$$

How to deal with imbalanced classes

Choose dice-loss over cross-entropy.

Balance your cross-entropy according to the class imbalance.

Extend the dice-loss to the Tverski loss.

Which weighs the influence
of the False Positives
and False Negative.

$$DL(\mathbf{p}, \hat{\mathbf{p}}) = \frac{\langle \mathbf{p}, \hat{\mathbf{p}} \rangle}{\langle \mathbf{p}, \hat{\mathbf{p}} \rangle + \beta \langle \mathbf{1} - \mathbf{p}, \hat{\mathbf{p}} \rangle + (1 - \beta) \langle \mathbf{p}, \mathbf{1} - \hat{\mathbf{p}} \rangle}$$

In our case $\beta = 1/8$

PASCAL Visual Object Classes

Pixel-wise segmentation of objects from a number of visual object classes in realistic scenes (i.e. not pre-segmented objects).

Annotations

Person, animals, vehicles, indoor.

Number of samples

6929 Pixel-wise
instance level annotations.

Metric

Mean Intersection over Union.

Image



Objects



<http://host.robots.ox.ac.uk/pascal/VOC/voc2012/index.html>

Neural Network Segmentation - Datasets and benchmarking

Common Objects in Context

COCO-Stuff augments 164K images with pixel-level stuff annotations for semantic segmentation.

Annotations

91 stuff classes (wall, grass, etc.) and 80 thing classes (person, elephant, etc.), as well as captions.

Number of samples

164000 dense pixel-level annotations and instance level annotations for things.

Metric

Mean Intersection over Union.



<https://github.com/nightrome/cocostuff>

MedakaHeart

Dataset to enable heart ventricle segmentation in Medaka fish for quantification of ventricular dimensions.

Annotations


Medaka heart ventricle.

Number of samples

1725 binary pixel-level annotations.

Metric

Dice coefficient, ventricular dimensions.



<https://osf.io/snb6p/>

Neural Network Segmentation - Datasets and benchmarking

Cityscapes

The Cityscapes Dataset focuses on semantic understanding of urban street scenes.

Annotations


City scene semantic and instance-wise pixel annotations (road, person, pole, etc.).

Number of samples

30 classes in 5000 fine and 20000 coarse annotated images.

Metric

Mean Intersection over Union
and Instance Intersection over Union.



<https://www.cityscapes-dataset.com/>

Neural Network Segmentation - Datasets and benchmarking

ISBI Segmentation Challenge


Goals is the automatic segmentation of neural structure.

Annotations

Transmission Electron Microscopy (ssTEM) of the Drosophila first instar larva ventral nerve cord (VNC).

Number of samples

30 consecutive images and the labelled boundary map.



http://brainiac2.mit.edu/isbi_challenge/

Metric

Rand Scoring (similarity measurement between partitions).

So how many samples do we need for segmentation?


So how many samples do we need for segmentation?,
Is an ill-posed question.

Neural Network Segmentation - Datasets and benchmarking

So how many samples do we need for segmentation?,
Is an ill-posed question.

Variance of the dataset

If your data is sequential or otherwise
very similar you will need more samples.



http://brainiac2.mit.edu/isbi_challenge/

Instances per sample
Complexity of the task


Neural Network Segmentation - Datasets and benchmarking

So how many samples do we need for segmentation?,
Is an ill-posed question.


Variance of the dataset

Instances per sample

Data augmentation approaches, such as cropping can harness the availability of multiple instances in one sample.



http://brainiac2.mit.edu/isbi_challenge/



<https://osf.io/snb6p/>

Complexity of the task



Neural Network Segmentation - Datasets and benchmarking

So how many samples do we need for segmentation?,
Is an ill-posed question.


Variance of the dataset
Instances per sample

Complexity of the task

The more complex the task, such as variance in shape, orientation, or intra sample variance, the more data is necessary to generalize well enough.



http://brainiac2.mit.edu/isbi_challenge/



<https://osf.io/snb6p/>

Architectures over time

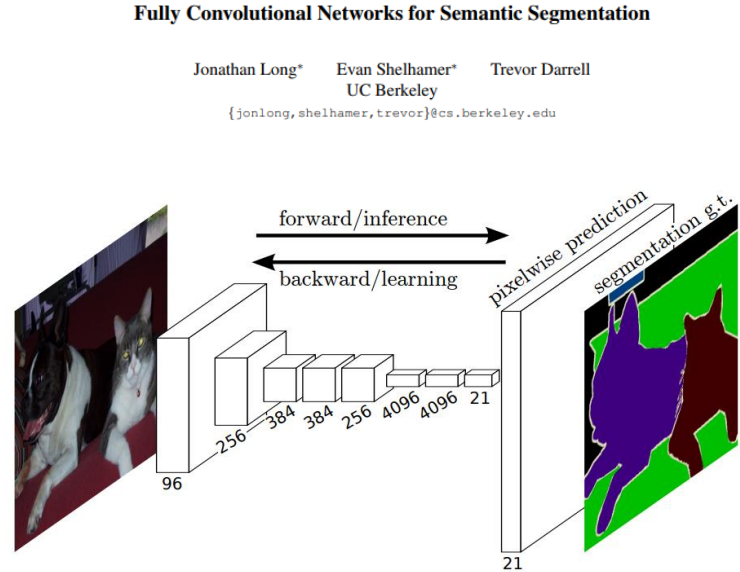
Fully Convolutional Network	2015
ParseNet	2015
Convolutional and Deconvolutional Networks	2015
U-Net	2015
Feature Pyramid Network	2016
Mask R-CNN	2017
DeepLab	2017

Fully Convolutional Network

First end-to-end trained Fully Convolutional Network for image segmentation.

Transfer Learning approach, modifying well known architectures (such as VGG16).

Ending with an upsampling layer with one channel per class.




https://people.eecs.berkeley.edu/~jonlong/long_shelhamer_fcn.pdf

TensorFlow Hub

A library *tensorflow_hub* for reusable machine learning modules in TensorFlow.

Such as text embedding, image feature vectors, and video classification.



<https://tfhub.dev/>

```


1 import tensorflow_hub as hub
2
3
4 def create_module_graph(module_spec):
5     """Creates a graph and loads Hub Module into it.
6     Args:
7         module_spec: the hub.ModuleSpec for the image module being used.
8     Returns:
9         graph: the tf.Graph that was created.
10        bottleneck_tensor: the bottleneck values output by the module.
11        resized_input_tensor: the input images, resized as expected by the module.
12        wants_quantization: a boolean, whether the module has been instrumented
13        with fake quantization ops.
14    """
15    height, width = hub.get_expected_image_size(module_spec)
16    with tf.Graph().as_default() as graph:
17        resized_input_tensor = tf.placeholder(tf.float32, [None, height, width, 3])
18        m = hub.Module(module_spec)
19        bottleneck_tensor = m(resized_input_tensor)
20        wants_quantization = any(node.op in FAKE_QUANT_OPS
21                                for node in graph.as_graph_def().node)
22    return graph, bottleneck_tensor, resized_input_tensor, wants_quantization
23
24
25

```

ParseNet

The global context of an image helps segmentation.

This is done by depicting the context in a global feature, which is later combined with the feature map.



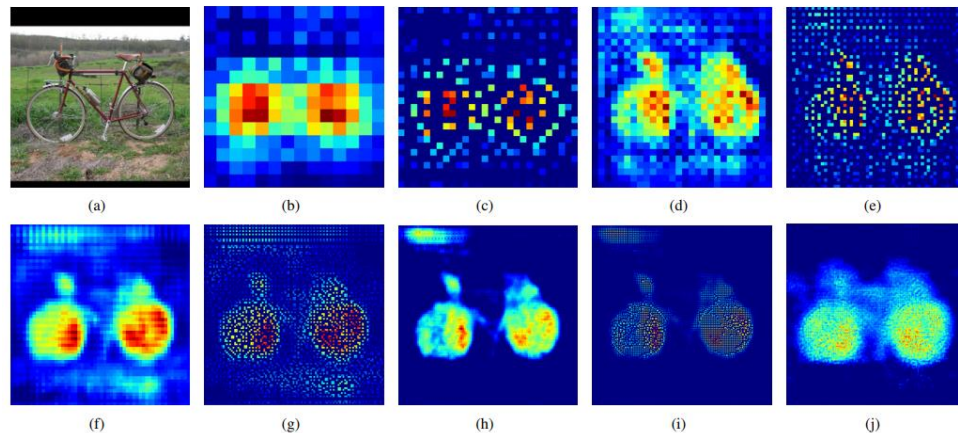

<https://arxiv.org/pdf/1506.04579.pdf>

Convolutional and Deconvolutional Networks

Introducing an encoder-decoder architecture.

From the convolutional encoding, the deconvolution branch generates a dense pixel-wise class probability map, by successive:

Unpooling, deconvolutions, and rectifications.




<https://arxiv.org/pdf/1505.04366.pdf>

Feature Pyramid Network

Introducing a bottom-up pathway to generate features at different scales.

Combined with a top-down pathway doing the upsampling, while infusing the high-level features through lateral connections.

On each stage of the pyramid a segmentation mask is predicted.




<https://arxiv.org/pdf/1612.03144.pdf>

Mask R-CNN

Using a Region Proposal Network to extract a manageable number of regions of interest.

A multi-task loss, combining losses for the bounding box coordinates, the class prediction and the segmentation mask improve the performance.

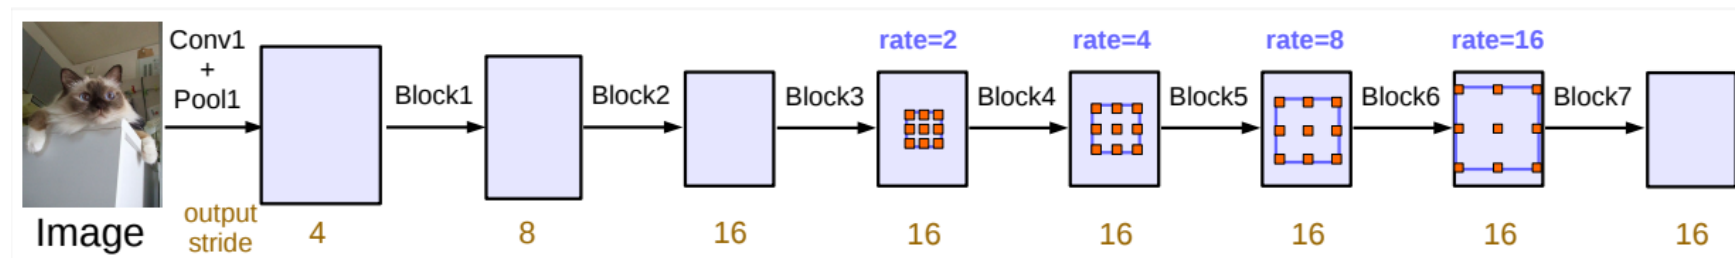


<https://arxiv.org/pdf/1703.06870.pdf>

DeepLabv3

Combining Atrous Convolutions (dilated convolutions) with a pyramidal architecture.

Atrous convolutions replace a combination of pooling, convolution and unpooling.



<https://arxiv.org/pdf/1706.05587.pdf>

Atrous Convolution


Introducing an additional parameter, called the dilation rate or rate.

Defining a spacing between the values in a filter map.

Atrous Convolution

Introducing an additional parameter, called the dilation rate or rate.

Defining a spacing between the values in a filter map.




Filter 3x3

Dilation rate: 2

Stride: 1

Padding: 1




Atrous Convolution

Introducing an additional parameter, called the dilation rate or rate.

Defining a spacing between the values in a filter map.

This enhances the field of view while keeping the computational cost low.




Filter 3x3

Dilation rate: 2

Stride: 1

Padding: 1



Machine Learning Methods for Automated Quantification of Ventricular Dimensions

Mark Schutera¹, Steffen Just², Jakob Gierten^{3, 4}, Ralf Mikut¹, Markus Reischl¹, Christian Pylatiuk^{1*}

1 Institute for Automation and Applied Informatics (IAI), Karlsruhe Institute of Technology (KIT),
Hermann-von-Helmholtz-Platz 1, 76344 Eggenstein, Germany

2 Department of Internal Medicine II, University of Ulm, Albert-Einstein-Allee 23, 89081 Ulm, Germany

3 Department of Pediatric Cardiology, University Hospital Heidelberg, Im Neuenheimer Feld 430, 69120 Heidelberg,
Germany

4 Centre for Organismal Studies Heidelberg, Heidelberg University, Im Neuenheimer Feld 230, 69120 Heidelberg,
Germany

Machine Learning Methods for Automated Quantification of Ventricular Dimensions

Project Description

Data Recording

Data Annotation

Data Augmentation

Deep Dive U-Net

Ventricular Dimensions

Results

Discussion

Project Description


Medaka fish contribute to the understanding of human cardiovascular diseases. In this context the quantification of cardiac functional parameters is fundamental.

Project Description

Medaka fish contribute to the understanding of human cardiovascular diseases.

In this context the quantification of cardiac functional parameters is fundamental.

Medaka prove to be a valuable whole organism-based in vivo model due to the visibility of the heart and other organs.



Project Description

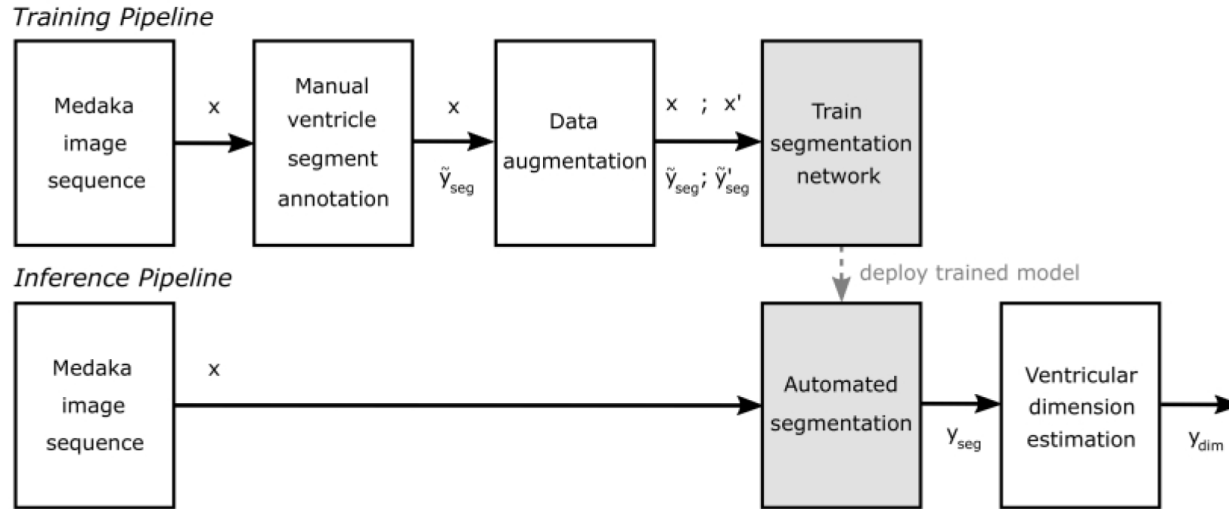
Medaka fish contribute to the understanding of human cardiovascular diseases. In this context the quantification of cardiac functional parameters is fundamental.

Medaka prove to be a valuable whole organism-based in vivo model due to the visibility of the heart and other organs.

Automated screening pipelines are thus invaluable for biomedical research.


Development of a pipeline for automated measurement and quantification of important cardiac functional parameters.

Project Description



Data Recording

Medaka larvae were imaged 1-2 days after hatching.



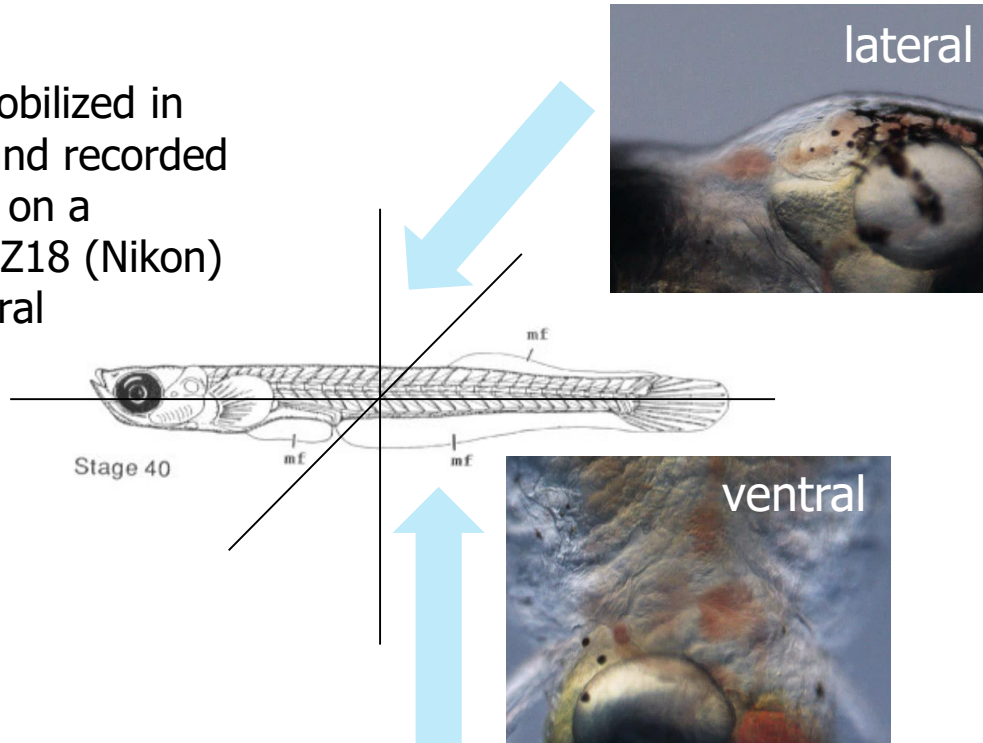
<http://mepd.cos.uni-heidelberg.de/mepd/forms/developmentalStages.jsf>

Medaka
image
sequence

Manual
ventricle
segment
annotation

Data Recording

Hatchlings were immobilized in 3% methylcellulose and recorded at room temperature on a stereomicroscope SMZ18 (Nikon) from ventral and lateral



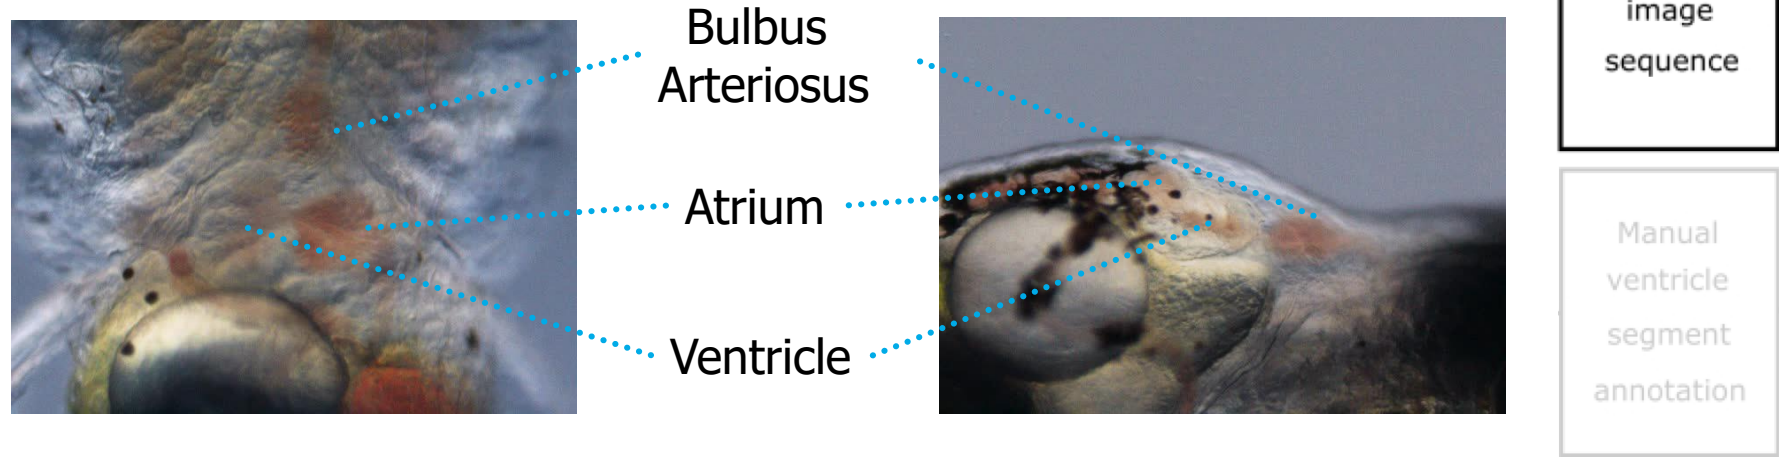
Medaka
image
sequence

Manual
ventricle
segment
annotation

<http://mepd.cos.uni-heidelberg.de/mepd/forms/developmentalStages.jsf>



Data Recording

At zoom magnification of 6x with 640x480 pixels (1 px = 1,115 μ m) and 15 frames per second (fps).



Data Recording

To reduce the computing effort, the images are converted to grayscale




Medaka
image
sequence

Manual
ventricle
segment
annotation

Data Recording

To reduce the computing effort, the images are converted to grayscale and rescaled to 256×256 px (.tif).



Medaka
image
sequence

Manual
ventricle
segment
annotation

Data Recording

The data is extracted from 63 videos in total comprising approx. 115 sec of image sequences or 1725 frames.



Medaka
image
sequence

Manual
ventricle
segment
annotation

Data Recording

The data is extracted from 63 videos in total comprising approx. 115 sec of image sequences.

Training and Validation Set

725 frames from 29 ventral sequences

500 frames from 20 lateral sequences



Medaka
image
sequence

Manual
ventricle
segment
annotation

Data Recording

The data is extracted from 63 videos in total comprising approx. 115 sec of image sequences.

Test Set Ventral

150 consecutive frames from a single sequence

5x25 consecutive frames from 5 sequences



Medaka
image
sequence

Manual
ventricle
segment
annotation

Data Recording

The data is extracted from 63 videos in total comprising approx. 115 sec of image sequences.

Training Set Lateral

150 consecutive frames from a single sequence
5x25 consecutive frames from 5 sequences




Medaka
image
sequence

Manual
ventricle
segment
annotation

Data Annotation

Each frame of the dataset has been annotated individually by manually labeling the present heart segment / ventricle, using the brush-tool of the pixel annotation tool v1.3.1.



<https://github.com/abreheret/PixelAnnotationTool>


Medaka
image
sequence

Manual
ventricle
segment
annotation

Data
augmentation

Data Annotation

The annotations are binary image masks (.tiff) whereas ground truth pixels are assigned to 1 and background pixels to 0.



<https://osf.io/snb6p/>

Medaka
image
sequence

Manual
ventricle
segment
annotation


Data
augmentation

Data Annotation

The annotations are binary image masks (.tif) whereas ground truth pixels are assigned to 1 and background pixels to 0.

Time devoted per frame
25 seconds

Time devoted for the dataset
14 hours



<https://osf.io/snb6p/>

Medaka
image
sequence


Manual
ventricle
segment
annotation

Data
augmentation

Data Augmentation

Confronted with the great effort required for data annotation, and the subsequent low amount of available ground truth (1225 frames),


data augmentation methods are applied.




Data Augmentation

The samples and the according ground truth are randomly augmented during training, by applying a set of augmentation methods.


horizontal flips



vertical flips



Zoom [0;0.1]



Manual
ventricle
segment
annotation


Data
augmentation

Train
segmentation
network


Data Augmentation

The samples and the according ground truth are randomly augmented during training, by applying a set of augmentation methods.



**height and
width shift [0;0.05]**



rotation [0;0.2]



shear [0;0.1]



Manual
ventricle
segment
annotation

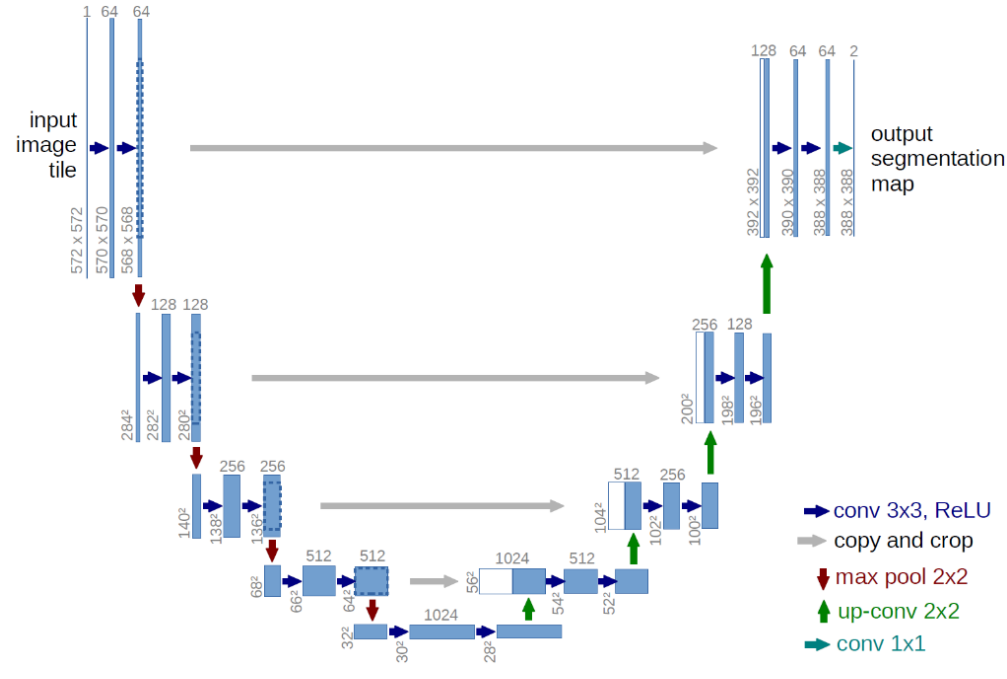
Data
augmentation

Train
segmentation
network


Neural Network Segmentation - Deep Dive U-Net

Deep Dive U-Net

The U-Net is a symmetric, deep convolutional neural network.




<https://lmb.informatik.uni-freiburg.de/people/ronneber/u-net/>



Neural Network Segmentation - Deep Dive U-Net

Deep Dive U-Net

The down sampling path computes high-level features with semantic information.



<https://lmb.informatik.uni-freiburg.de/people/ronneber/u-net/>

Data
augmentation


Train
segmentation
network

Ventricular dimension estimation

Neural Network Segmentation - Deep Dive U-Net

Deep Dive U-Net

The up sampling path computes spatially localized patterns in the image.



<https://lmb.informatik.uni-freiburg.de/people/ronneber/u-net/>

Data
augmentation

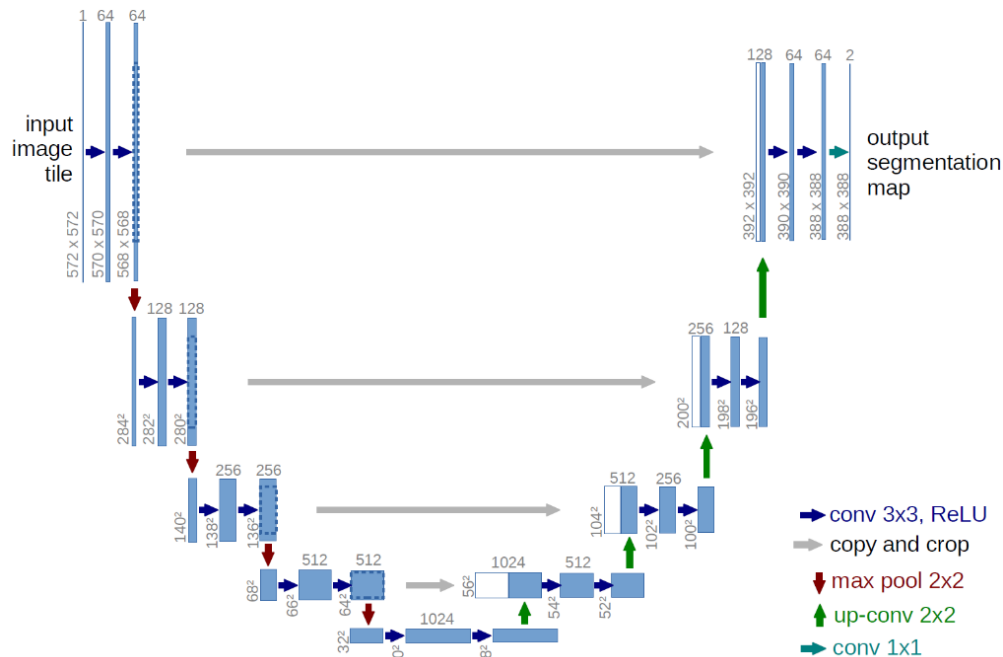
Train
segmentation
network

Ventricular
dimension
estimation

Deep Dive U-Net

Both paths are brought together by skip connections.

Combining semantic patterns with localization information.



<https://lmb.informatik.uni-freiburg.de/people/ronneber/u-net/>

Data
augmentation

Train
segmentation
network

Ventricular
dimension
estimation

Training Parameters

Adam optimizer with a learning rate of $\epsilon = 1e^{-5}$

Batch size of 8

Training schedule 60 epochs

300 iterations per epoch

Data augmentation

Batch normalization in the convolution layers

Dropout in the bottleneck layers

Data
augmentation

Train
segmentation
network

Ventricular
dimension
estimation


Dice loss

Similar to the IoU (Intersection over union), and thus easy to interpret.

+1 is a smoothing factor for numeric stability.

Robust with respect to imbalanced classes.

$$dice\ loss = 1 - \frac{2 \sum(\tilde{\mathbf{y}} \odot \mathbf{y}_{seg}) + 1}{\sum(\tilde{\mathbf{y}}^2) + \sum(\mathbf{y}_{seg}^2) + 1}.$$



Ventricular dimensions

Heart rate

Ventricular volumes

Fractional shortenings


are linked to the **cardiovascular cycle**.

Train
segmentation
network

Ventricular
dimension
estimation

Neural Network Segmentation - Ventricular dimensions

To estimate the ventricular dimensions, the frame-based predicted segments have to be **transferred into the temporal domain**.



temporal domain

Train
segmentation
network

Ventricular
dimension
estimation

Neural Network Segmentation - Ventricular dimensions

To estimate the ventricular dimensions, the frame-based predicted segments have to be transferred into the temporal domain.

And we need to determine the cardiovascular cycle frequency.

For example through a feature such as the **segment area a in px over time**

$$a^{(t)} = \sum_{u=0}^{image\ width} \sum_{v=0}^{image\ height} (y_{seg}(u, v))^{(t)}$$

Train
segmentation
network

Ventricular
dimension
estimation

Neural Network Segmentation - Ventricular dimensions

To estimate the ventricular dimensions, the frame-based predicted segments have to be transferred into the temporal domain.

And we need to determine the cardiovascular cycle frequency.

For example through a feature such as the **segment area a in px over time**

$$a^{(t)} = \sum_{u=0}^{image\ width} \sum_{v=0}^{image\ height} (y_{seg}(u, v))^{(t)}$$

Other features could be minor axis, major axis, equivalent diameter, and so on

Train
segmentation
network

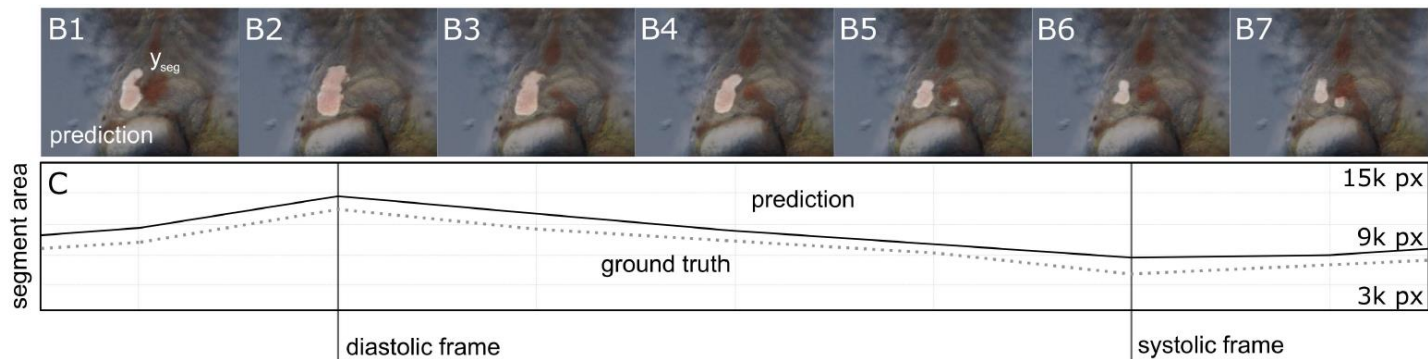
Ventricular
dimension
estimation

Neural Network Segmentation - Ventricular dimensions

To estimate the ventricular dimensions, the frame-based predicted segments have to be transferred into the temporal domain.

And we need to determine the cardiovascular cycle frequency.

And a sliding window approach for **segment area peak detection**




Train
segmentation
network

Ventricular
dimension
estimation

What do the peaks mean?

Diastolic frame: Maximum heart relaxation with maximal heart volume

Systolic frame: Maximum heart tension with minimal heart volume



Train
segmentation
network

Ventricular
dimension
estimation

Heart rate

The heart rate (HR) is consequently estimated from, **offset** o the amount of frames between two systolic or diastolic frames, which correspond to one heartbeat.

The beats per frame are subsequently transformed into the time space by compensating for the sample frequency $f = 15\text{fps}$.

$$HR = \frac{f}{o}$$

Train
segmentation
network

Ventricular
dimension
estimation

Ventricular volumes
Fractional shortenings

Neural Network Segmentation - Ventricular dimensions

Heart rate

Ventricular volumes



Train
segmentation
network

Ventricular
dimension
estimation

are approximated by a prolate spheroidal shape, based on the major and minor axis of the predicted heart segment. The minor axis e_{minor} and major axis e_{major} serve as semi-diameter:

$$V = \frac{1}{6} \pi e_{minor}^2 e_{major}$$

Fractional shortenings

Neural Network Segmentation - Ventricular dimensions

Heart rate

Ventricular volumes

Fractional shortenings

is the fraction of a diastolic dimension (D_{dia}) that is lost in systole (D_{sys}).

The most important fractional shortening is the **shortening of the ventricles volume**, or ejection fraction is thus given by the percentage of volume change between a diastolic frame and its subsequent systolic frame.

$$V = \frac{V_{dia} - V_{sys}}{V_{dia}}$$

Train
segmentation
network

Ventricular
dimension
estimation

Results

Automated heart segmentation

Detection rate for diastolic / systolic frame pairs is 100%.

Mean error 0.53 frames and max error is 1.0 frames.

(Ground truth is in brackets)

Current heart beat	diastolic frame; systolic frame	heart rate [bps]	minor fractional shortening [%]	major fractional shortening [%]
1.	2; 6 (2; 7)	1.50 (1.50)	41.33 (47.29)	52.18 (38.80)
2.	11; 15 (11; 16)	1.67 (1.67)	40.87 (30.90)	57.58 (38.69)
3.	19; 23 (19; 24)	1.58 (1.58)	28.57 (48.92)	53.46 (34.57)
4.	28; 31 (28; 32)	1.46 (1.67)	29.37 (43.21)	47.16 (28.68)
5.	36; 42 (36; 40)	1.69 (1.50)	42.56 (45.43)	60.73 (41.05)
6.	45; 49 (45; 49)	1.52 (1.50)	39.36 (42.56)	58.66 (31.75)
7.	53; 59 (54; 58)	1.82 (1.67)	33.81 (45.71)	61.26 (34.37)
8.	62; 65 (62; 66)	1.58 (1.52)	38.61 (47.74)	58.32 (49.29)
9.	70; 74 (70; 76)	1.52 (1.58)	50.33 (43.99)	53.06 (37.22)
10.	78; 84 (79; 84)	1.58 (1.50)	44.15 (49.85)	55.30 (31.20)
11.	87; 92 (88; 93)	1.58 (1.67)	39.93 (45.48)	43.54 (39.53)
12.	96; 100 (96; 101)	1.58 (1.67)	44.27 (47.76)	58.00 (44.15)
13.	104; 109 (104; 109)	1.58 (1.58)	47.96 (42.90)	59.27 (41.57)
14.	112; 118 (113; 117)	1.58 (1.58)	38.71 (46.30)	64.93 (20.96)
15.	121; 126 (122; 125)	1.58 (1.52)	41.75 (44.24)	55.37 (22.83)
16.	130; 134 (130; 135)	1.58 (1.58)	20.39 (52.80)	59.46 (37.15)
17.	138; 143 (139; 143)	1.67 (1.52)	43.92 (48.11)	52.80 (37.61)
18.	146; 151 (147; 153)	1.50 (1.50)	35.01 (36.93)	56.59 (35.89)
<i>mean</i>		1.59 (1.57)	38.93 (45.01)	55.98 (35.85)
<i>std</i>		0.08 (0.07)	7.10 (4.65)	4.98 (6.62)
<i>median</i>		1.59 (1.58)	38.93 (45.59)	55.98 (37.19)
<i>lower quartile</i>		1.53 (1.50)	7.10 (43.40)	4.98 (32.41)
<i>upper quartile</i>		1.58 (1.65)	7.10 (47.75)	4.98 (39.35)

Results

Automated heart segmentation

Detection rate 100%

Max detection error is 1 frame

Ventricular dimension estimation is robust over multiple cardiovascular cycles.

Current heart beat	diastolic frame; systolic frame	heart rate [b/s]	minor fractional shortening [%]	major fractional shortening [%]
1.	2; 6 (2; 7)	1.50 (1.50)	41.33 (47.29)	52.18 (38.80)
2.	11; 15 (11; 16)	1.67 (1.67)	40.87 (30.90)	57.58 (38.69)
3.	19; 23 (19; 24)	1.58 (1.58)	28.57 (48.92)	53.46 (34.57)
4.	28; 31 (28; 32)	1.46 (1.67)	29.37 (43.21)	47.16 (28.68)
5.	36; 42 (36; 40)	1.69 (1.50)	42.56 (45.43)	60.73 (41.05)
6.	45; 49 (45; 49)	1.52 (1.50)	39.36 (42.56)	58.66 (31.75)
7.	53; 59 (54; 58)	1.82 (1.67)	33.81 (45.71)	61.26 (34.37)
8.	62; 65 (62; 66)	1.58 (1.52)	38.61 (47.74)	58.32 (49.29)
9.	70; 74 (70; 76)	1.52 (1.58)	50.33 (43.99)	53.06 (37.22)
10.	78; 84 (79; 84)	1.58 (1.50)	44.15 (49.85)	55.30 (31.20)
11.	87; 92 (88; 93)	1.58 (1.67)	39.93 (45.48)	43.54 (39.53)
12.	96; 100 (96; 101)	1.58 (1.67)	44.27 (47.76)	58.00 (44.15)
13.	104; 109 (104; 109)	1.58 (1.58)	47.96 (42.90)	59.27 (41.57)
14.	112; 118 (113; 117)	1.58 (1.58)	38.71 (46.30)	64.93 (20.96)
15.	121; 126 (122; 125)	1.58 (1.52)	41.75 (44.24)	55.37 (22.83)
16.	130; 134 (130; 135)	1.58 (1.58)	20.39 (52.80)	59.46 (37.15)
17.	138; 143 (139; 143)	1.67 (1.52)	43.92 (48.11)	52.80 (37.61)
18.	146; 151 (147; 153)	1.50 (1.50)	35.01 (36.93)	56.59 (35.89)
<i>mean</i>		1.59 (1.57)	38.93 (45.01)	55.98 (35.85)
<i>std</i>		0.08 (0.07)	7.10 (4.65)	4.98 (6.62)
<i>median</i>		1.59 (1.58)	38.93 (45.39)	55.98 (37.19)
<i>lower quartile</i>		1.53 (1.50)	7.10 (43.40)	4.98 (32.41)
<i>upper quartile</i>		1.58 (1.65)	7.10 (47.75)	4.98 (39.35)

Results

Automated heart segmentation

Detection rate 100%

Max detection error is 1 frame

Ventricular dimension estimation is robust over multiple cardiovascular cycles.

Reported human performance is between 5% to 12% on fractional shortenings.

Current heart beat	diastolic frame; systolic frame	heart rate [b/s]	minor fractional shortening [%]	major fractional shortening [%]	
1.	2; 6	(2; 7)	1.50 (1.50)	41.33 (47.29)	52.18 (38.80)
2.	11; 15	(11; 16)	1.67 (1.67)	40.87 (30.90)	57.58 (38.69)
3.	19; 23	(19; 24)	1.58 (1.58)	28.57 (48.92)	53.46 (34.57)
4.	28; 31	(28; 32)	1.46 (1.67)	29.37 (43.21)	47.16 (28.68)
5.	36; 42	(36; 40)	1.69 (1.50)	42.56 (45.43)	60.73 (41.05)
6.	45; 49	(45; 49)	1.52 (1.50)	39.36 (42.56)	58.66 (31.75)
7.	53; 59	(54; 58)	1.82 (1.67)	33.81 (45.71)	61.26 (34.37)
8.	62; 65	(62; 66)	1.58 (1.52)	38.61 (47.74)	58.32 (49.29)
9.	70; 74	(70; 76)	1.52 (1.58)	50.33 (43.99)	53.06 (37.22)
10.	78; 84	(79; 84)	1.58 (1.50)	44.15 (49.85)	55.30 (31.20)
11.	87; 92	(88; 93)	1.58 (1.67)	39.93 (45.48)	43.54 (39.53)
12.	96; 100	(96; 101)	1.58 (1.67)	44.27 (47.76)	58.00 (44.15)
13.	104; 109	(104; 109)	1.58 (1.58)	47.96 (42.90)	59.27 (41.57)
14.	112; 118	(113; 117)	1.58 (1.58)	38.71 (46.30)	64.93 (20.96)
15.	121; 126	(122; 125)	1.58 (1.52)	41.75 (44.24)	55.37 (22.83)
16.	130; 134	(130; 135)	1.58 (1.58)	20.39 (52.80)	59.46 (37.15)
17.	138; 143	(139; 143)	1.67 (1.52)	43.92 (48.11)	52.80 (37.61)
18.	146; 151	(147; 153)	1.50 (1.50)	35.01 (36.93)	56.59 (35.89)
mean		1.59 (1.57)	38.93 (45.01)	55.98 (35.85)	
std		0.08 (0.07)	7.10 (4.65)	4.98 (6.62)	
median		1.59 (1.58)	38.93 (45.39)	55.98 (37.19)	
lower quartile		1.53 (1.50)	7.10 (43.40)	4.98 (32.41)	
upper quartile		1.58 (1.65)	7.10 (47.75)	4.98 (39.35)	

Results

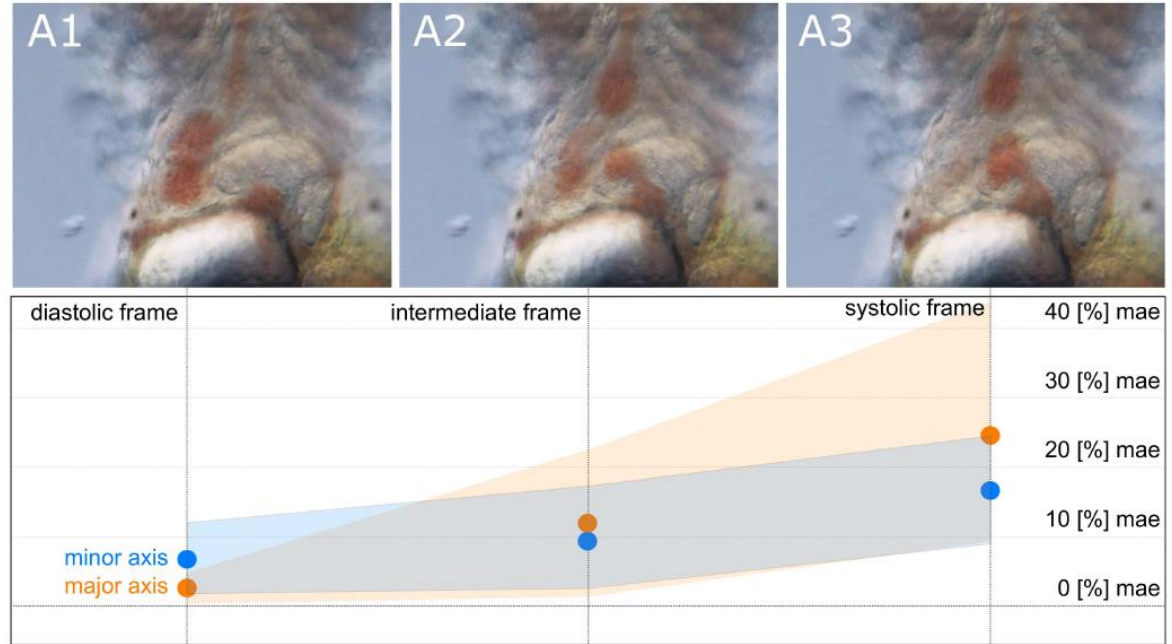
Automated heart segmentation

Detection rate 100%

Max detection error is 1 frame





Ventricular dimension
estimation is robust

Errors are largest for systolic
frames, this is when humans
also fail to label precisely.



Results

Data-driven end-to-end automated heart segmentation from ventral view and lateral view.
Robust estimation of ventricular dimensions.



Codebase

<https://osf.io/snb6p/>

*Estimation of further
ventricular dimensions seems
feasible*

Want to get your hands dirty?
Let me know.

Journal Paper
Tutorial

Machine Learning Methods for Automated Quantification of Ventricular Dimensions

Contributors: [Mark Schutera](#), Steffen Just, Jakob Gierten, Ralf Mikut, Markus Reischl, Christian Pylatiuk

Date created: 2019-03-24 02:53 AM | Last Updated: 2019-05-29 11:47 PM

[Create DOI](#)


Category:  Project

Description:

Machine learning methods for automated quantification of ventricular dimensions (for further details see Wiki bellow or take a closer look on our paper).

License: GNU General Public License (GPL) 3.0

Wiki



See demonstration of our algorithm and framework on the test set data:

https://youtu.be/ISbX_XbwXq0

The medaka (*Oryzias latipes*) and the zebrafish (*Danio rerio*) are used as a model organism for a variety of subjects in biomedical research the here presented work aims to study the potential of automated ventricular dimension estimation through heart segmentation in medaka. For more on this, it's t...

[Read More](#)

Neural Network Segmentation - Take it from here

Codebase
Journal Paper

Tutorial Kudos to Hendrik Vogt

jupyter UNet-TUT Last Checkpoint: Last Monday at 3:30 PM (unsaved changes)

File Edit View Insert Cell Kernel Help Not Trusted Python 2

Segmentation Code: <https://github.com/zhouhaotunet> Nette Anwendungsbeispiel für das Tutorial (Kannst ja mal schauen ob das auf der CPU konvergiert, batch_size auf 1) <https://tutorials.mepirical.com/image-segmentation-with-unet/>

U-net Segmentation with KERAS

Introduction

This Tutorial will face up with the U-net Segmentation, which was presented in this [paper](#). Firstly the theory behind the architecture will be discussed to subsequently implement this neural network within this Jupyter Notebook, to segment the cell boundaries from a gray image, see Fig 1.

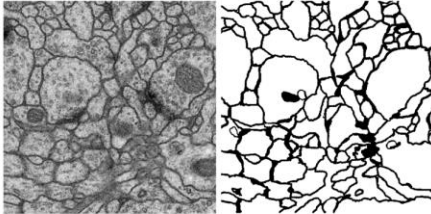



Figure 1: Right: input image | Left: Prediction of the cell boundaries

Motivation

While convolutional neural networks often need a very large amount of training data, it is preferable to use the training data which might be not that large, more efficiently. Furthermore the neural nets often just classify a whole image instead of given any further specific information about the localization of its containing objects or scene. Hence with the U-Net each pixel of the image is labeled using the information of its surrounding pixels. This is primarily done with the known convolution and maxpooling, but to avoid a loss of localization information you get with these functions, the U-Net architecture was conceived. The U-Net architecture also eliminate the two drawbacks of a former approaches which where the high redundancy, due to the overlapping tiles (will be explained below) and the loss of context when minimize the patches.

What is U-Net



References

- [6] Lex Fridman. MIT 6.S094: Deep Learning for Self-Driving Cars <https://selfdrivingcars.mit.edu/>, 2019. Zugriff: 20.06.2019.
- [7] ^ Sørensen, T. (1948). "A method of establishing groups of equal amplitude in plant sociology based on similarity of species and its application to analyses of the vegetation on Danish commons". *Kongelige Danske Videnskabernes Selskab*. 5 (4): 1–34.
- [8] A guide to convolution arithmetic for deep learning. V. Dumoulin, and F. Visin. (2016) arxiv:1603.07285.
- [9] Deep learning. Yann LeCun, Yoshua Bengio & Geoffrey Hinton Nature volume521, pages436–444 (2015)
- [10] K. Simonyan and A. Zisserman. Very deep convolutional networks for large-scale image recognition. *CoRR*, abs/1409.1556, 2014. 1, 2, 3, 5
- [11] A. Bréhéret. Pixel annotation tool. *GitHub Repository*, 2017.
- [12] O. Ronneberger, P. Fischer, and T. Brox. U-net: Convolutional networks for biomedical image segmentation. *Medical Image Computing and Computer-Assisted Intervention*, 9351:234–241, 2015.
- [13] T. Hoage, Y. Ding, and X. Xu. Quantifying cardiac functions in embryonic and adult zebrafish. *Cardiovascular Development*, 843:11–20, 2012.

Thanks for your time

Questions?

Contact!

mark.schutera@kit.edu

