



Applied Deep Learning

Deep Reinforcement Learning

Mark Schutera

Karlsruhe Institute of Technology | ZF Friedrichshafen AG | Hochschule Ravensburg-Weingarten

1. Deep Learning Foundations
-
3. Transfer Learning and Object Detection
-
5. Segmentation Networks
-
- 8. Deep Reinforcement Learning**
-
10. Generative Adversarial Networks
-
12. Recurrent Neural Networks

Course overview

Project proposal

One page on introduction, methods, dataset

Deadline 3. Lecture

Intermediate presentation

Ten minutes on achievements, problems, next steps

Due 7. Lecture

Final presentation

Screencast (Slides and Audio)

Due 13. Lecture (20.01.2020)

Final documentation

Documentation and code on github

Deadline 13. Lecture (20.01.2020)

Course features

Sli.do

Every question matters.

Get the app.

Ask questions (with slide number)
or vote on other students' questions
during the lecture.

And give direct feedback.

#TOBEDETERMINED

Questions will be covered
immediately or in the next lecture in
more depth.

Github

Find slides, tutorials, flashcards and
references on Github.

[https://github.com/schutera/
DeepLearningLecture](https://github.com/schutera/DeepLearningLecture) **Schutera**

You found typos, additional
material such as links, algorithms,
papers, literature or want to
contribute to the slides and lecture
notes..

..Feel free to contribute, e-mail me.

Course features

Sli.do

Every question matters.

Get the app.

Ask questions (with slide number)
or vote on other students' questions
during the lecture.

#TOBEDETERMINED

Questions will be covered directly or
in the next lecture in more depth.

Github

Typos, additional material such as
links, algorithms, paper, literature,
lecture notes..

..Feel free to contribute.

Grade Bonus .3

Prepare flashcards based on Ian
Goodfellow's Deep Learning Book

- Commit to flashcard set by
emailing me, first come first serve
- Must be comprehensive

Introduction to reinforcement learning

About reinforcement learning

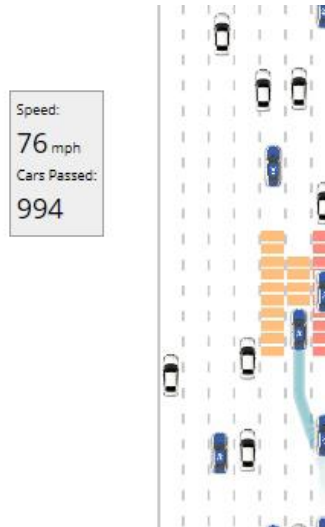
Reinforcement learning a problem statement

Reinforcement learning framework

Reinforcement Learning with neural networks


Introduction – Reinforcement learning

Reinforcement learning (RL) is an area of machine learning concerned with how software agents ought to take actions in an environment so as to maximize some notion of cumulative reward.



Introduction – Reinforcement learning

Learn to make good sequences of decisions



What makes reinforcement learning different from other machine learning paradigms

- Feedback / **supervision** directly comes from the **environment**
- **Model interacts** with the **environment**
or at least its own state in the environment

- **Dense highway traffic**

- **80** mph max. Speed
- **7** Lanes
- **20** Vehicles
- **1 - 11** Trainable Vehicles

- **Deep Reinforcement Competition**

The screenshot displays the DeepTraffic web application. On the left, a vertical highway simulation shows multiple lanes with cars. A status box indicates 'Speed: 40 mph' and 'Cars Passed: 136'. Below the simulation, there are controls for 'Road Overlay' (set to 'None') and 'Simulation Speed' (set to 'Normal'). A 'Load Custom Image' button and a 'Teal' color selector are also present. In the center, a car with the ZF logo is shown. To the right, the 'DeepTraffic' title is followed by links to 'Main Page', 'Leaderboard', and 'About DeepTraffic'. A text block states: 'Americans spend 8 billion hours stuck in traffic every year. Deep neural networks can help!'. Below this is a code editor with a JavaScript snippet for a reinforcement learning environment. Further down are buttons for 'Apply Code/Reset Net', 'Save Code/Net to File', 'Load Code/Net from File', and 'Submit Model to Competition'. A line graph shows a performance metric over time. At the bottom right, there are buttons for 'Run Training' and 'Start Evaluation Run', followed by a diagram of a 'Value Function Approximating Neural Network' with layers: input(19), fc(1), relu(1), fc(5), and regression(5). A 'Vehicle Skins' link is also visible.

```
1  
2 //<![CDATA[  
3  
4 // a few things don't have var in front of them - they update already  
   existing variables the game needs  
5 lanesSide = 0;  
6 patchesAhead = 1;  
7 patchesBehind = 0;  
8 trainIterations = 10000;
```

<https://selfdrivingcars.mit.edu/deeptraffic/>

Deep Reinforcement Learning with the Deep Q-Network

Deep Reinforcement Learning Decision Making

- **Environment**

Either real-world or simulation of some sort

- Action
- State
- Reward



Deep Reinforcement Learning with the Deep Q-Network


Deep Reinforcement Learning Decision Making

- Environment

- **Action a**

Agent's actions affect
the subsequent data it receives

- State
- Reward



Deep Reinforcement Learning with the Deep Q-Network

Deep Reinforcement Learning


Decision Making

- Environment
- Action

- **State s**

Feedback is delayed, not instantaneous

- Reward




Deep Reinforcement Learning with the Deep Q-Network

Deep Reinforcement Learning Decision Making

- Environment
- Action
- State

- **Reward r**


There is no supervisor, only a reward signal



Deep Reinforcement Learning

Time really matters (sequential signal stream)

- Environment
- Action
- State
- Reward




Maze Example

Conventional Reinforcement Learning

- **Environment**

Maze, Start and End Position


- Action
- State
- Reward



Maze Example

Conventional Reinforcement Learning

- Environment
- **Action**
Move left, right, up, down
- State
- Reward



Maze Example


Conventional Reinforcement Learning

- Environment
- Action

- **State**

Position in one of the maze cells


- Reward



Maze Example

Conventional Reinforcement Learning


- Environment
- Action
- State
- **Reward**
-1 for valid move, +1 for transition into the end position



Maze Example

Conventional Reinforcement Learning

- Environment
- Action
- State
- **Reward**
-1 for valid move, +1 for transition into the end position



What happens for a reward of 0 for a valid move?

Introduction to reinforcement learning

Reinforcement Learning with neural networks

Q-Learning

Fixed target network

Experience replay

Reward clipping

Frame skipping

Reward function

Self-Play

Deep dive Deep-Q-Learning

Datasets and benchmarking and current success stories

Reinforcement Learning and why it does not work yet

Q-Learning

Value-based algorithm

Lookup table of values with one entry for every state-action pair

$$Q(s, a)$$

Deep Q-Learning

Representation Learning

Lookup table of values with one entry for every state-action pair

$$Q(s, a)$$

High dimensional state-action spaces, make a mere Q-value function inapplicable

In this case a parametrized value function (neural network) is needed

$$Q(s, a; \theta_i)$$

Deep Q-Learning

Representation Learning

Idea follows a trial-and-error strategy, exploring the states and iteratively updating the state-action values

$$Q(s, a)$$

All actions ought to be **repeatedly sampled** in **all states** to ensure sufficient exploration.

Deep Q-Learning

Representation Learning

Idea follows a trial-and-error strategy,
exploring the states and iteratively updating
the state-action values

$$Q(s, a)$$

How does our final / optimal policy look like?

Given a state calculate the expected maximal
reward of an action, following our policy.

$$Q^*(s, a) = \max_{\pi} \mathbb{E} \left[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \mid s_t = s, a_t = a, \pi \right]$$

Deep Q-Learning

Loss function

$$L_i(\theta_i) = \mathbb{E} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Our network is trained the same way we are used to – by backpropagating a loss and adjusting our weights.

Deep Q-Learning

Loss function

$$L_i(\theta_i) = \mathbb{E} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Our network is trained the same way we are used to – by backpropagating a loss and adjusting our weights.

The central idea here is not to find the best action, but to learn an accurate estimation of the Q-value function.

Deep Q-Learning

Loss function

$$L_i(\theta_i) = \mathbb{E} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Our network is trained the same way we are used to – by backpropagating a loss and adjusting our weights.

The central idea here is not to find the best action, but to learn an accurate estimation of the Q-value function.

The loss is thus the difference between the true **reward** from the environment

And the **estimated reward** of our model

Deep Q-Learning

Fixed Target Network

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

θ_i^-

The parameters for the target network are only updated every C iterations

Deep Q-Learning

Fixed Target Network

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

The parameters for the target network are only updated every C iterations

This reduces the risk of oscillations or divergence and increases stability of the training process

Deep Q-Learning

Experience Replay

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

State action pairs together with the expected reward and resulting state transition are stored for N time steps.

Deep Q-Learning

Experience Replay

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

State action pairs together with the expected reward and resulting state transition are stored for N time steps.

During training we then draw from the experience replay memory

Deep Reinforcement Learning with Neural Networks – Tricks that make you converge

Deep Q-Learning

Experience Replay

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

State action pairs together with the expected reward and resulting state transition are stored for N time steps.

During training we then draw from the experience replay memory

Actions are drawn randomly with probability
Epsilon ϵ

ϵ -greedy policy

And trained within a mini-batch.

Deep Q-Learning

ϵ -greedy policy

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Exploration

$\epsilon = 1$

Meaning actions are completely randomly selected resulting in maximum exploration within our action-state space

Exploitation

Deep Q-Learning

ϵ -greedy policy

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Exploration

Exploitation

$\epsilon = 0$

Meaning no random actions what so ever.

Actions are selected with respect to the maximum expected reward

Deep Q-Learning

ϵ -greedy policy

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Exploration

Exploitation

During training which ϵ -value would you want to use?

Deep Q-Learning

ϵ -greedy policy

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Exploration

Exploitation

During training which ϵ -value would you want to use?

$\epsilon = 1$ for **full exploration first** (think baby)

This at some point might be inefficient since you do not profit from already learned action-state pairs.

Decrease ϵ over time (think grown-up)

Deep Q-Learning

ϵ -greedy policy

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Exploration

Exploitation

During inference which ϵ -value would you want to use?

Deep Q-Learning

ϵ -greedy policy

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Exploration

Exploitation

During inference which ϵ -value would you want to use?

$\epsilon = 0$ for **full exploitation** (think adult who has his routines and does not really reflect what she is doing)

This at some point might get you stuck in a deadlock

Thus **allow for some randomness** in your choice of action $\epsilon = 0.05$ (think adult who tries to drink water instead of soda)

Deep Q-Learning

Reward Clipping

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Keep the range of values of a reward within 1 and -1

Deep Q-Learning

Reward Clipping

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Keep the range of values of a reward within 1 and -1

This scaling allows for robust weight updates, keeping the **error derivatives** small

And also allows for using the **same learning** over multiple different environments

Deep Q-Learning

Frame Skipping

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

As Reinforcement Learning has a **sequential nature**, it is to be assumed that **adjacent frames are similar** and contribute similar information gain.

Deep Q-Learning

Frame Skipping

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

As Reinforcement Learning has a sequential nature, it is to be assumed that adjacent frames are similar and contribute similar information gain.

Idea is to **reduce sample rate** or in other words to skip frames during training to increase information gain and thus training and **sample and compute efficiency**

Deep Reinforcement Learning with Neural Networks – Tricks that make you converge

Deep Q-Learning

Reward Function Definition

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Definition of the reward function is the means to influence what's learned.

This might often not end up as expected



Deep Q-Learning

Reward Function Definition

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Shaped rewards

Increasing rewards in states that are closer to the end goal.

Easier to learn, but also prone to induce bias.

Sparse rewards

Reward shaping

Deep Q-Learning

Reward Function Definition

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Shaped rewards

Sparse rewards

Reward at the goal state.

Hard to learn and the lack of positive reinforcement might even make it too difficult.

Rewards shaping

Deep Q-Learning

Reward Function Definition

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

Shaped rewards

Sparse rewards

Reward shaping

Design an objective function that diverts from the actual objective.

This can help in settings with sparse rewards.

Deep Reinforcement Learning with Neural Networks – Tricks that make you converge

Deep Q-Learning

Reward Function Definition

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

What kind of reward function is this?

How could you minimize the introduced bias?

How could you transform it into a differently shaped reward?



Deep Q-Learning


Self-Play

Idea is that agent plays itself and thus improves on the learning objective. There are variations to this strategy

Both sides learn simultaneously

Left side learns

Right side learns



<https://gym.openai.com/envs/Pong-v0/>

References

- [1] [Bergstra and Bengio, 2012] James Bergstra and Yoshua Bengio. Random search for hyper-parameter optimization. *Journal of Machine Learning Research*, 13(Feb):281–305, 2012.
- [2] [Fridman *et al.*, 2018] Lex Fridman, Benedikt Jenik, and Jack Terwilliger. Deeptraffic: Driving fast through dense traffic with deep reinforcement learning. *CoRR*, abs/1801.02805, 2018.
- [3] [Mnih *et al.*, 2015] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A Rusu, Joel Veness, Marc G Bellemare, Alex Graves, Martin Riedmiller, Andreas K Fidjeland, Georg Ostrovski, et al. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529, 2015.
- [4] [Silver *et al.*, 2017] David Silver, Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Matthew Lai, Arthur Guez, Marc Lanctot, Laurent Sifre, Dhharshan Kumaran, Thore Graepel, et al. Mastering chess and shogi by self-play with a general reinforcement learning algorithm. *arXiv preprint arXiv:1712.01815*, 2017.
- [5] [Tuyls and Weiss, 2012] Karl Tuyls and Gerhard Weiss. Multiagent learning: Basics, challenges, and prospects. *Association for the Advancement of Artificial Intelligence*, 2012.
- [6] [Olivas *et al.*, 2009] Emilio Soria Olivas, Jose David Martin Guerrero, Marcelino Martinez Sober, Jose Rafael Magdalena Benedito, and Antonio Jose Serrano Lopez. *Handbook Of Research On Machine Learning Applications and Trends: Algorithms, Methods and Techniques - 2 Volumes*. Information Science Reference - Imprint of: IGI Publishing, Hershey, PA, 2009.

OpenAI Gym – Classic Control Problems

Several classic control problems that long have been used to evaluate reinforcement learning algorithms

Description

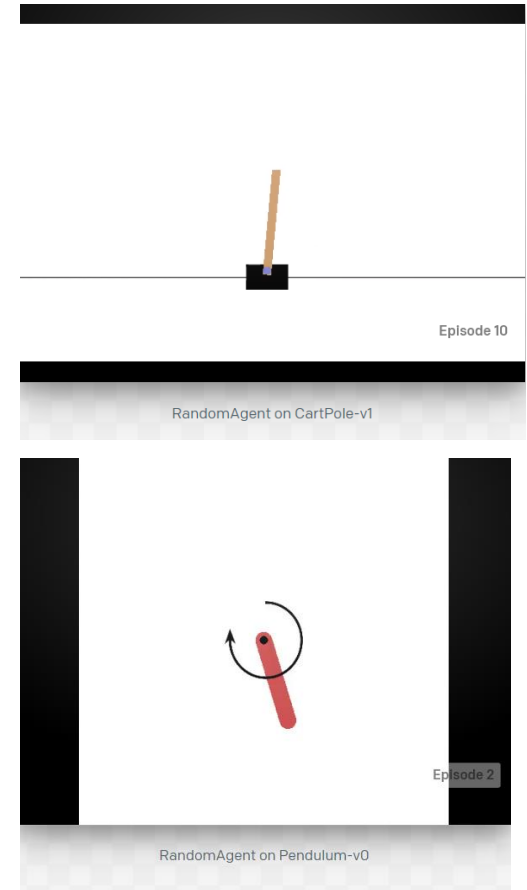
Toolkit for developing and comparing reinforcement learning algorithms

Environments

CartPole

Pendulum

MountainCar



<http://gym.openai.com/>

OpenAI Gym – Atari Games

Several classic control problems that long have been used to evaluate reinforcement learning algorithms

Description

Large state spaces, and or large action spaces. Long planning horizons. Different variants of the same game allows for transfer learning or generalization evaluations

Environments

SpaceInvaders

Pong

Breakout



[Breakout-ram-v0](#)

Maximize score in the game
Breakout, with RAM as input



[SpaceInvaders-ram-v0](#)

Maximize score in the game
SpaceInvaders, with RAM
as input

<http://gym.openai.com/>

OpenAI Gym – Continuous Control Systems

Continuous control problems in locomotion tasks

Description

Large state spaces, and or large action spaces. Long planning horizons. Different variants of the same game allows for transfer learning or generalization evaluations


Environments

Ant

Half Cheetah

Hopper

Humanoid



Ant-v2

Make a 3D four-legged robot walk.



Humanoid-v2

Make a 3D two-legged robot walk.

<http://gym.openai.com/>

Starcraft Gym – PySC2

It exposes Blizzard Entertainment's StarCraft II Machine Learning API as a Python RL Environment. This is a collaboration between DeepMind and Blizzard to develop StarCraft II into a rich environment for RL research.

Description

Mini-Map games and Full-Map Leaderboards with extensive evaluation tooling



<http://starcraftgym.com/>

AlphaStar Versus Serral

<https://www.youtube.com/watch?v=DMXvkbAtHNY>

Best practice

Benchmarking metrics and experimental guidelines

Number of trials

Results might vary significantly by just changing the random seed. Thus we need to statistically ground our experiments

Hyperparameter Tuning

Benchmark Environments and Metrics

Best practice

Benchmarking metrics and experimental guidelines

Number of trials

Hyperparameter Tuning

Ensuring a fair comparison between learning algorithms. For example by ablation analysis (removing a certain hyperparameter, such as no dropout)

Benchmark Environments and Metrics

Best practice

Benchmarking metrics and experimental guidelines

Number of trials

Hyperparameter Tuning

Benchmark Environments and Metrics

Ideally a large mixture of environments should be covered. Same holds for metrics, do not only report on the optimal metrics, report on all.

Success over time

Backgamon agent based on Reinforcement Learning (<i>Tesauro et al.</i>)	1995
Superhuman performance in Atari games (<i>Mnih et al.</i>)	2013
Mastering the game of Go (<i>Silver et al.</i>)	2016
Mastering the game of StarCraft II (<i>Deep Mind</i>)	2019

Field of Application

Robotics, Autonomous Driving, etc.



<https://deepmind.com/blog/article/alphastar-mastering-real-time-strategy-game-starcraft-ii>

TD-gammon

Backgammon playing agent entirely based on reinforcement learning and self-play.

Based on value-function, similar to Q-learning.
Approximation of the value-function based on multi-layer perceptron with one hidden layer.

Temporal Difference Learning and TD-Gammon

By Gerald Tesauro




Figure 1. An illustration of the multilayer perceptron architecture used in TD-Gammon's neural network. This architecture is also used in the popular backpropagation learning procedure. Figure reproduced from [9].

<https://cling.csd.uwo.ca/cs346a/extra/tdgammon.pdf>

Deep Q-Network

Deep Q-Network with experience replay, fixed target network, mini-batch training, and convolution architecture.

Exploration and exploitation set by epsilon variable.

Ability to master Atari 2600 computer games using pixel input.

Playing Atari with Deep Reinforcement Learning

Volodymyr Mnih Koray Kavukcuoglu David Silver Alex Graves Ioannis Antonoglou

Daan Wierstra Martin Riedmiller

DeepMind Technologies

{vlad,koray,david,alex.graves,ioannis,daan,martin.riedmiller} @ deepmind.com

	B. Rider	Breakout	Enduro	Pong	Q*bert	Seaquest	S. Invaders
Random	354	1.2	0	-20.4	157	110	179
Sarsa [3]	996	5.2	129	-19	614	665	271
Contingency [4]	1743	6	159	-17	960	723	268
DQN	4092	168	470	20	1952	1705	581
Human	7456	31	368	-3	18900	28010	3690
HNeat Best [8]	3616	52	106	19	1800	920	1720
HNeat Pixel [8]	1332	4	91	-16	1325	800	1145
DQN Best	5184	225	661	21	4500	1740	1075

Table 1: The upper table compares average total reward for various learning methods by running an ϵ -greedy policy with $\epsilon = 0.05$ for a fixed number of steps. The lower table reports results of the single best performing episode for HNeat and DQN. HNeat produces deterministic policies that always get the same score while DQN used an ϵ -greedy policy with $\epsilon = 0.05$.

<https://arxiv.org/pdf/1312.5602.pdf>

Deep Q-Network

One system that is able to learn several games without any tweaking from game to game.

Reward clipping

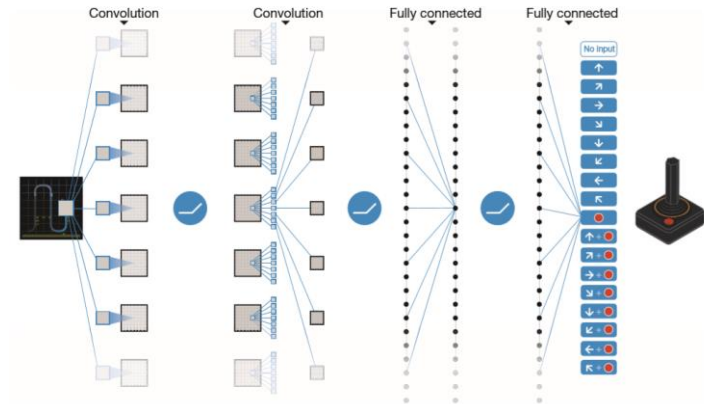
High-dimensional state spaces and high dimensional action spaces

LETTER

doi:10.1038/nature14236

Human-level control through deep reinforcement learning

Volodymyr Mnih^{1*}, Koray Kavukcuoglu^{1*}, David Silver^{1*}, Andrei A. Rusu¹, Joel Veness¹, Marc G. Bellemare¹, Alex Graves¹, Martin Riedmiller¹, Andreas K. Fiedjeland¹, Georg Ostrovski¹, Stig Petersen¹, Charles Beattie¹, Amir Sadik¹, Ioannis Antonoglou¹, Helen King¹, Dhharshan Kumaran¹, Daan Wierstra¹, Shane Legg¹ & Demis Hassabis¹



<https://storage.googleapis.com/deepmind-media/dqn/DQNNaturePaper.pdf>

Rainbow

Extensions to the Deep Q-Network (ablation study)

Double Q-Learning

Prioritized Replay

Dueling Networks

Multi-step learning

Distributional Reinforcement Learning

Noisy Nets

Rainbow: Combining Improvements in Deep Reinforcement Learning

Matteo Hessel
DeepMind

Joseph Modayil
DeepMind

Hado van Hasselt
DeepMind

Tom Schaul
DeepMind

Georg Ostrovski
DeepMind

Will Dabney
DeepMind

Dan Horgan
DeepMind

Bilal Piot
DeepMind

Mohammad Azar
DeepMind

David Silver
DeepMind




Figure 1: **Median human-normalized performance** across 57 Atari games. We compare our integrated agent (rainbow-colored) to DQN (grey) and six published baselines. Note that we match DQN's best performance after 7M frames,

<https://arxiv.org/pdf/1710.02298.pdf>

Alpha Star

Combination of Deep Reinforcement Learning with an LSTM module, pointer network

And a novel multi-agent, population-based learning algorithm.

Each agent experiences 200 years of real-time StarCraft play

StarCraft II: A New Challenge for Reinforcement Learning

Oriol Vinyals Timo Ewalds Sergey Bartunov Petko Georgiev
Alexander Sasha Vezhnevets Michelle Yeo Alireza Makhzani Heinrich Küttler
John Agapiou Julian Schrittwieser John Quan Stephen Gaffney Stig Petersen
Karen Simonyan Tom Schaul Hado van Hasselt David Silver Timothy Lillicrap
DeepMind

Kevin Calderone Paul Keet Anthony Brunasso David Lawrence
Anders Ekermo Jacob Repp Rodney Tsing
Blizzard



A VISUALISATION OF THE ALPHASTAR AGENT DURING GAME TWO OF THE MATCH AGAINST MANA. THIS SHOWS THE GAME FROM THE AGENT'S POINT OF VIEW. THE RAW OBSERVATION INPUT TO THE NEURAL NETWORK, THE NEURAL NETWORK'S INTERNAL ACTIVATIONS, SOME OF THE CONSIDERED ACTIONS THE AGENT CAN TAKE SUCH AS WHERE TO CLICK AND WHAT TO BUILD, AND THE PREDICTED OUTCOME. MANA'S VIEW OF THE GAME IS ALSO SHOWN, ALTHOUGH THIS IS NOT ACCESSIBLE TO THE AGENT.

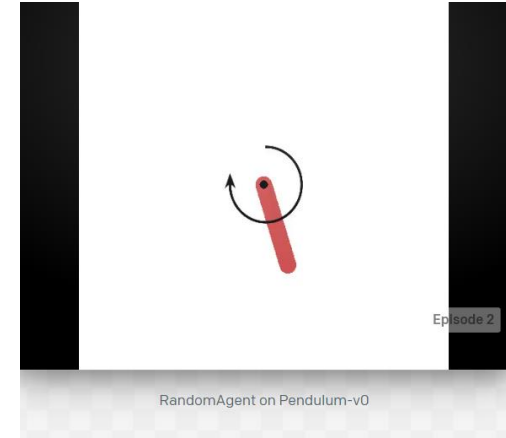
<https://deepmind.com/blog/article/alphastar-mastering-real-time-strategy-game-starcraft-ii>

Deep Reinforcement Learning is nice but ...

... real-world applications are still missing for the most part

... it learns to model environments that are unknown, that is usually not a requirement in domain specific problems

Technical Mechanics and Control Theory



<http://gym.openai.com/>

Deep Reinforcement Learning is nice but ...

... real-world applications are still missing for the most part

... it learns to model environments that are unknown, that is usually not a requirement in domain specific problems

Technical Mechanics and Control Theory



ATLAS Boston Dynamics

Paper

Deploying Rainbow to benchmark on current Traffic Light Control policies. Applied to real-world intersections in Friedrichshafen.

Tutorial

Distributed traffic light control at uncoupled intersections with real-world topology by deep reinforcement learning

Mark Schutera

Institute for Automation and Applied Informatics
Karlsruhe Institute of Technology
Research and Development
ZF Friedrichshafen AG
mark.schutera@kit.edu

Niklas Goby

Chair for Information Systems Research
University of Freiburg
IT Innovation Chapter Data Science
ZF Friedrichshafen AG
niklas.goby@is.uni-freiburg.de

Stefan Smolarek

IT Innovation Chapter Data Science
ZF Friedrichshafen AG
stefan.smolarek@zff.com

Markus Reischl

Institute for Automation and Applied Informatics
Karlsruhe Institute of Technology
markus.reischl@kit.edu

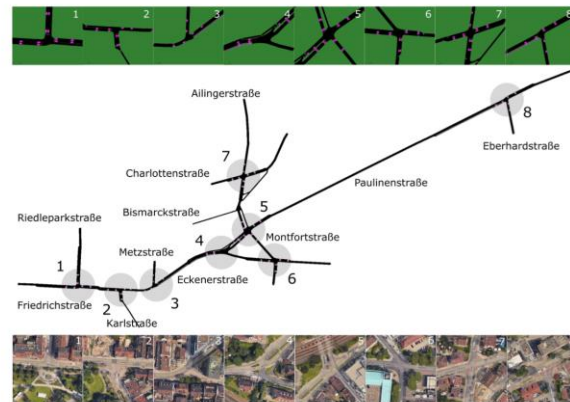


Figure 1: Overview of the Friedrichshafen roadnetwork with streetnames and the locations of the considered junctions in the center. In the top row the junctions are displayed as being present in SUMO. On the bottom row the google maps visualizations of the intersections themselves are shown.

<https://arxiv.org/pdf/1811.11233.pdf>

Paper

Analysis and comparison of transfer learning with multi-agent learning in highway traffic.

Extended analysis of congestion and vehicle behavior.

Tutorial

Transfer Learning versus Multi-agent Learning regarding Distributed Decision-Making in Highway Traffic

Mark Schutera^{1,4}, Niklas Goby^{2,3}, Dirk Neumann², Markus Reischl¹

¹ Institute for Automation and Applied Informatics, Karlsruhe Institute of Technology

² Chair for Information Systems Research, University of Freiburg

³ IT Innovation Chapter Data Science, ZF Friedrichshafen AG

⁴ Research and Development, ZF Friedrichshafen AG




Figure 1: Two screenshots from the micro-traffic simulation. The highlighted cells depict the catchment areas of the safety system, which automatically slows down the car to prevent collisions. The vehicles with the logo represent trainable agents, while those without a logo are not trainable and exhibit random behavior. The left figure shows the training process of a core network ANN_{core} , whereas on the right figure illustrates the pretrained core network being deployed among multiple agents.

<https://arxiv.org/pdf/1810.08515.pdf>

Paper Tutorial



Thanks for your time

Questions?

Contact!

mark.schutera@kit.edu

